

The Graphics Pipeline

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)

Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

Input:

Geometric model:

Description of all object, surface, and
light source geometry and transformations

Lighting model:

Computational description of object and
light properties, interaction (reflection)

Synthetic Viewpoint (or Camera):

Eye position and viewing frustum

Raster Viewport:

Pixel grid onto which image plane is mapped

Output:

Colors/Intensities suitable for framebuffer display
(For example, 24-bit RGB value at each pixel)

Modeling Transformations

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- 3D models defined in their own coordinate system (object space)
- Modeling transforms orient the models within a common coordinate frame (world space)



Object space

World space

Illumination (Shading) (Lighting)

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)


Visibility / Display

- Vertices lit (shaded) according to material properties, surface properties (normal) and light sources
- Local lighting model (Diffuse, Ambient, Phong, etc.)



Lighting Simulation

- Direct illumination
 - Ray casting
 - Polygon shading
- Global illumination
 - Ray tracing
 - Monte Carlo methods
 - Radiosity methods



Viewing Transformation

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Maps world space to eye space
- Viewing position is transformed to origin & direction is oriented along some axis (usually z)



Clipping

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)


Scan Conversion
(Rasterization)

Visibility / Display

- Transform to Normalized Device Coordinates (NDC)




- Portions of the object outside the view volume (view frustum) are removed




Why clip?

- We don't want to waste time rendering objects that are outside the **viewing window** (or clipping window)





What is clipping?

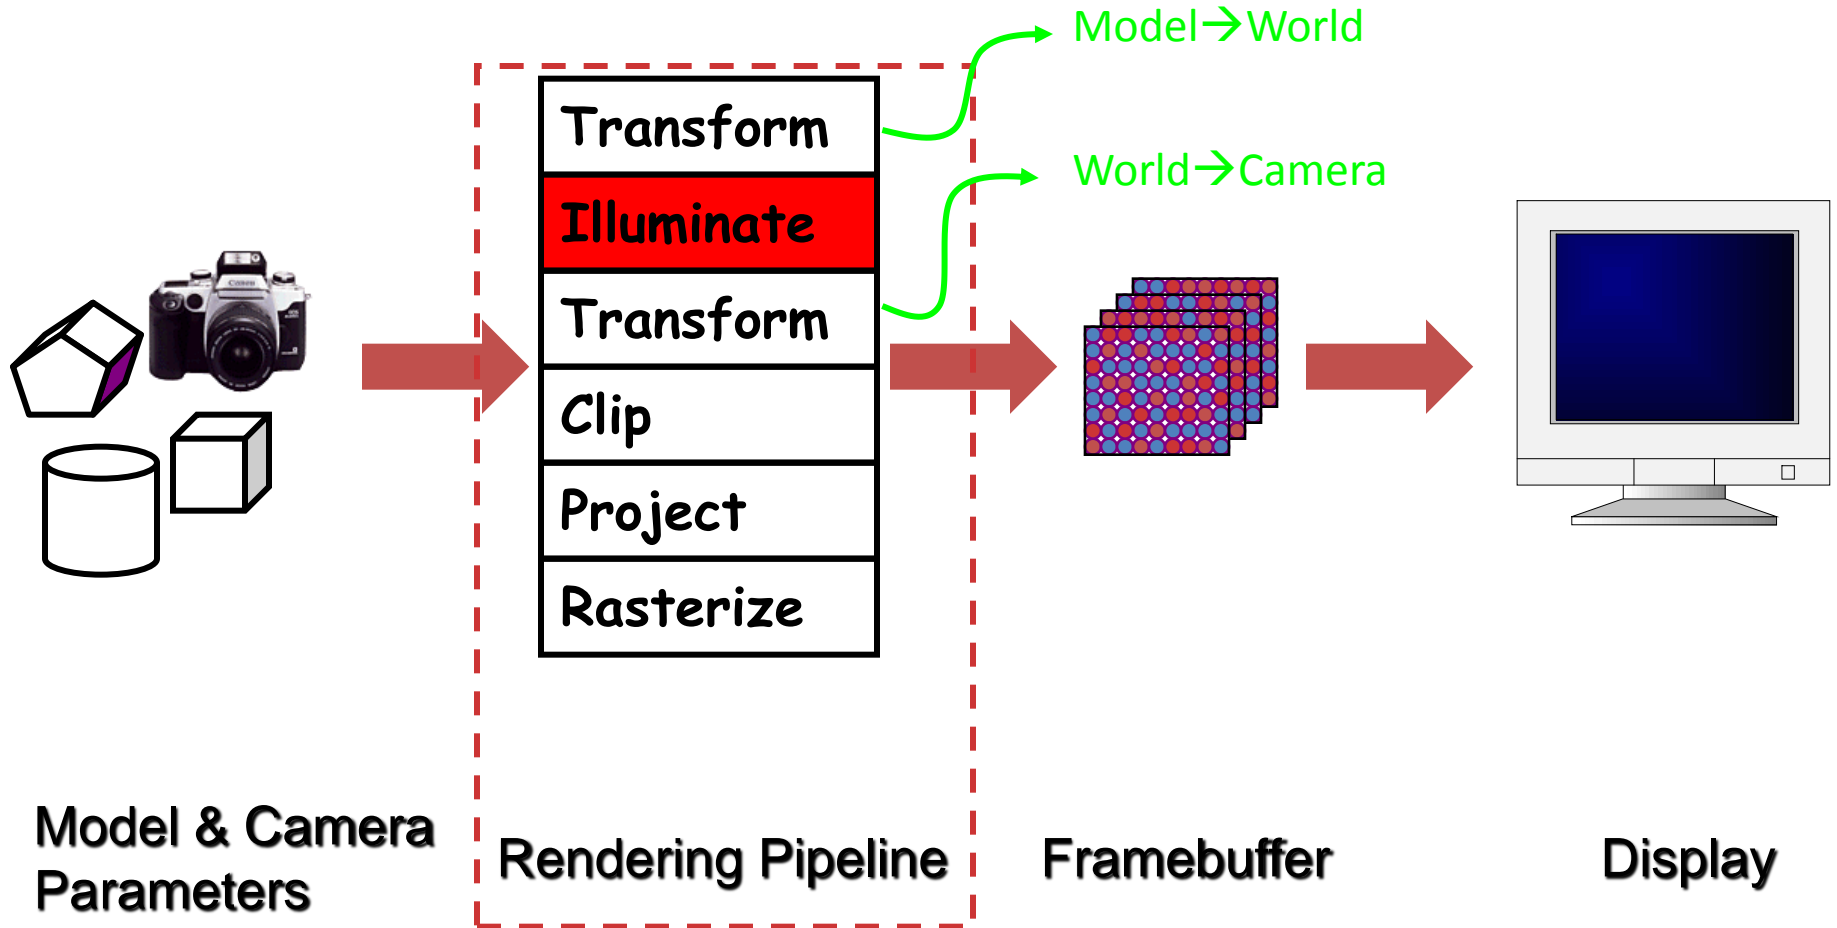
- Analytically calculating the portions of primitives within the view window



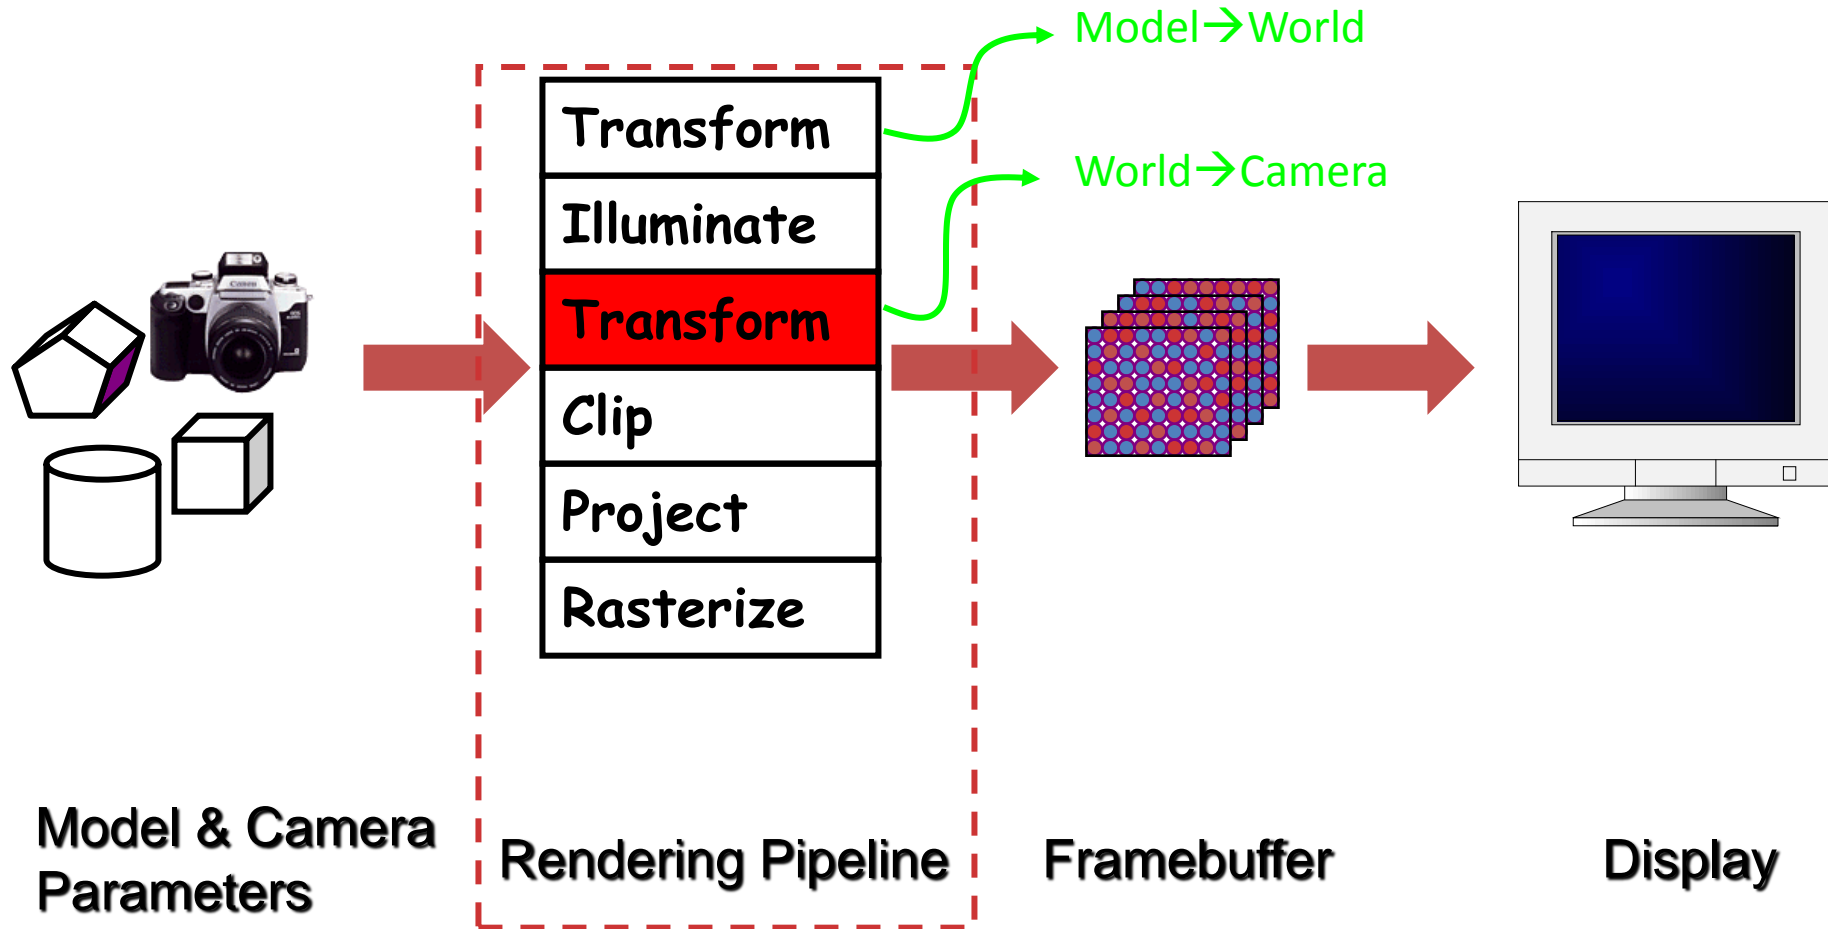
Clip to what?





Why illuminate before clipping?



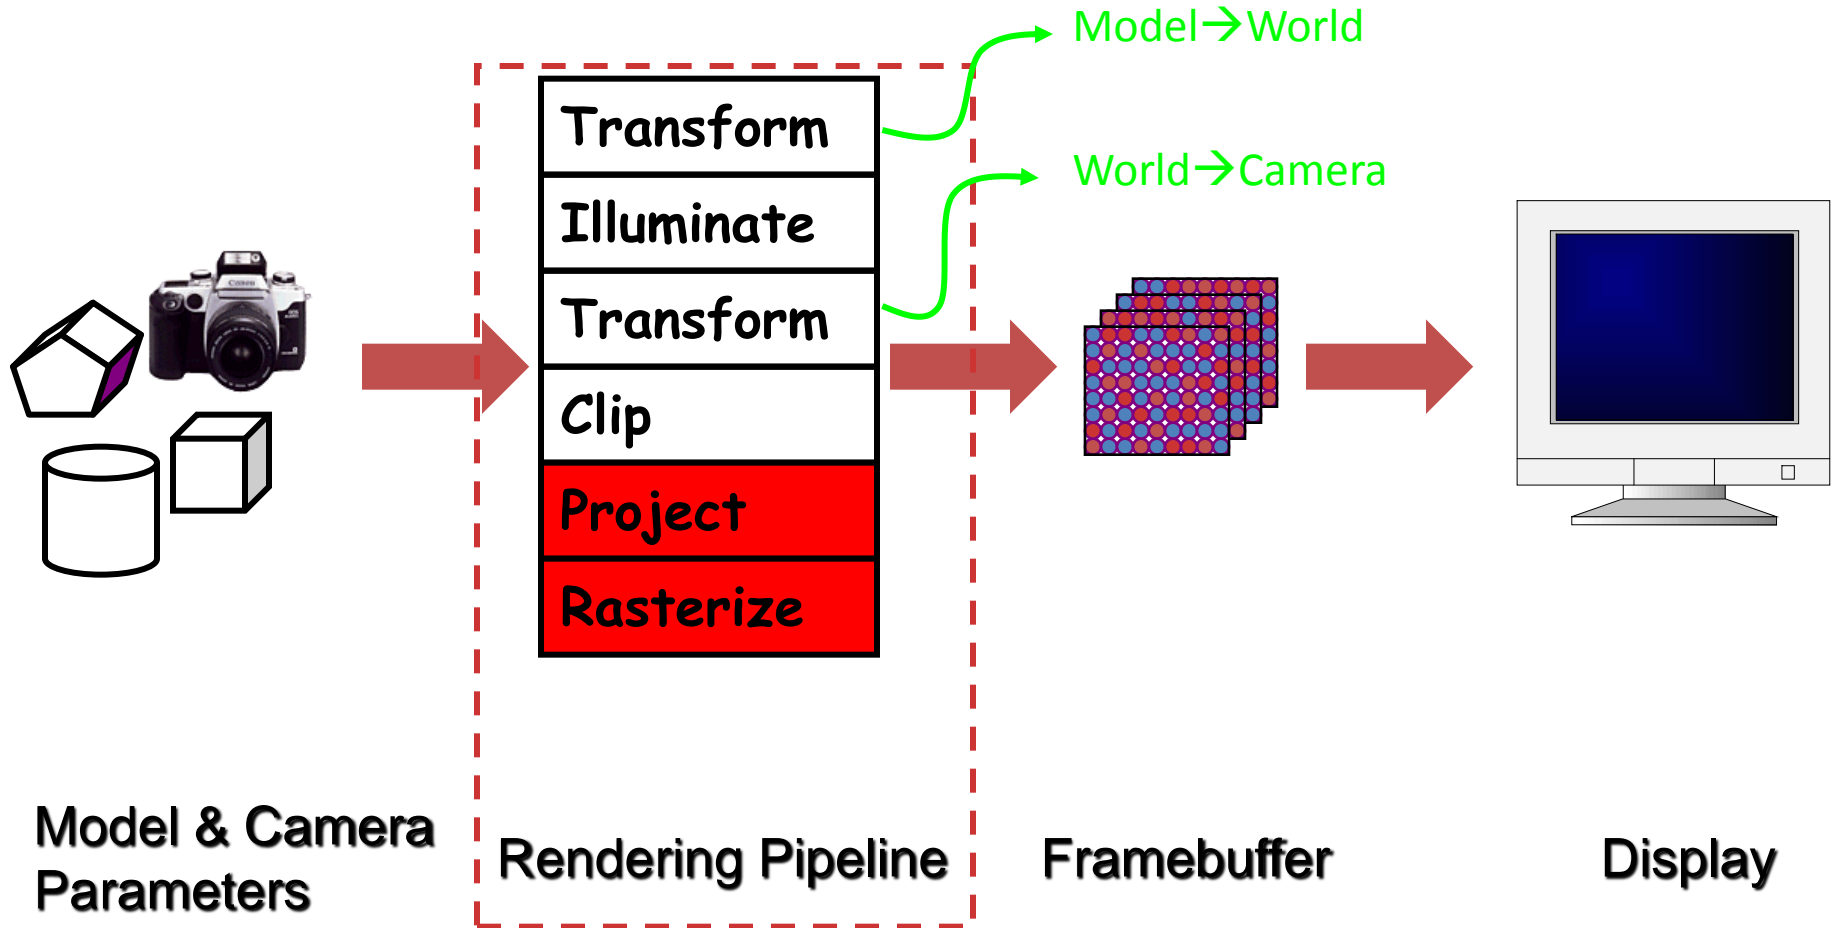
Why World \rightarrow Camera before clipping?



Clip to what?



Remind me why I care again




Why Clip?

- Bad idea to rasterize outside of framebuffer bounds
- Also, don't waste time scan converting pixels outside window


Clipping

- The naïve approach to clipping lines:
 for each line segment
 for each edge of view_window
 find intersection point
 pick “nearest” point
 if anything is left, draw it
- What do we mean by “nearest”?
- How can we optimize this?




Trivial Accepts

- Big optimization: trivial accept/rejects
- How can we quickly determine whether a line segment is entirely inside the view window?
- A: test both endpoints.




Trivial Rejects

- How can we know a line is outside view window?
- A: if both endpoints on wrong side of same edge, can trivially reject line




Clipping Lines To Viewport

- Combining trivial accepts/rejects
 - Trivially **accept** lines with both endpoints **inside all edges of the view window**
 - Trivially **reject** lines with both endpoints **outside the same edge of the view window**
 - Otherwise, reduce to trivial cases **by splitting into two segments**



Cohen-Sutherland Line Clipping

- Divide **view window** into regions defined by **window edges**
- Assign each region a 4-bit **outcode**:
 - Bit 1 indicates y-value of points are above y_{\max}



Cohen-Sutherland Line Clipping


- For each line segment
 - Assign an outcode to each vertex
 - If both outcodes = 0, trivial accept
 - Same as performing `if (bitwise OR = 0)`
 - Else
 - bitwise *AND* vertex outcodes together
 - if result $\neq 0$, trivial reject

Cohen-Sutherland Line Clipping

- If line cannot be trivially accepted or rejected, subdivide so that one or both segments can be discarded
 - Pick an edge of view window that the line crosses (*how?*)
 - Intersect line with edge (*how?*)
 - Discard portion on wrong side of edge and assign new outcode to new vertex
 - Apply trivial accept/reject tests; repeat if necessary


Cohen-Sutherland Line Clipping

- If line cannot be trivially accepted or rejected, subdivide so that one or both segments can be discarded
- Pick an edge that the line crosses
 - Check against edges in same order each time
 - For example: top, bottom, right, left




Cohen-Sutherland Line Clipping

- Intersect line with edge (how?)




Cohen-Sutherland Line Clipping

- Discard portion on wrong side of edge and assign outcode to new vertex
- Apply trivial accept/reject tests and repeat if necessary



Cohen-Sutherland Line Clipping

- Discard portion on wrong side of edge and assign outcode to new vertex
- Apply trivial accept/reject tests and repeat if necessary



View Window Intersection Code

- $(x_1, y_1), (x_2, y_2)$ intersect with vertical edge at x_{right}
 - $y_{\text{intersect}} = y_1 + m(x_{\text{right}} - x_1)$
 - where $m = (y_2 - y_1) / (x_2 - x_1)$
- $(x_1, y_1), (x_2, y_2)$ intersect with horizontal edge at y_{bottom}
 - $x_{\text{intersect}} = x_1 + (y_{\text{bottom}} - y_1) / m$
 - where $m = (y_2 - y_1) / (x_2 - x_1)$

Cohen-Sutherland Review

- Use opcodes to quickly eliminate/include lines
 - Best algorithm when trivial accepts/rejects are common
- Must compute viewing window clipping of remaining lines
 - Non-trivial clipping cost
 - Redundant clipping of some lines
- More efficient algorithms exist

Solving Simultaneous Equations

- Equation of a line
 - Slope-intercept (explicit equation): $y = mx + b$
 - Implicit Equation: $Ax + By + C = 0$
 - Parametric Equation: Line defined by two points, P_0 and P_1
 - $\mathbf{P}(t) = \mathbf{P}_0 + (\mathbf{P}_1 - \mathbf{P}_0) t$, where \mathbf{P} is a vector $[x, y]^T$
 - $x(t) = x_0 + (x_1 - x_0) t$
 - $y(t) = y_0 + (y_1 - y_0) t$

Parametric Line Equation

- Describes a finite line
- Works with vertical lines (like the viewport edge)
 - $0 \leq t \leq 1$
 - Defines line between P_0 and P_1
 - $t < 0$
 - Defines line before P_0
 - $t > 1$
 - Defines line after P_1

Parametric Lines and Clipping

- Define each line in parametric form:
 - $P_0(t) \dots P_{n-1}(t)$
- Define each edge of view window in parametric form:
 - $P_L(t), P_R(t), P_T(t), P_B(t)$
- Perform Cohen-Sutherland intersection tests using appropriate view window edge and line

Line / Edge Clipping Equations

- Faster line clippers use parametric equations

- Line 0:

- $x^0 = x_0^0 + (x_1^0 - x_0^0) t^0$

- $y^0 = y_0^0 + (y_1^0 - y_0^0) t^0$

View Window Edge L:

- $x^L = x_0^L + (x_1^L - x_0^L) t^L$

- $y^L = y_0^L + (y_1^L - y_0^L) t^L$

- $x_0^0 + (x_1^0 - x_0^0) t^0 = x_0^L + (x_1^L - x_0^L) t^L$

- $y_0^0 + (y_1^0 - y_0^0) t^0 = y_0^L + (y_1^L - y_0^L) t^L$


- Solve for t^0 and/or t^L

Cyrus-Beck Algorithm

- We wish to optimize line/line intersection
 - Start with parametric equation of line:
 - $P(t) = P_0 + (P_1 - P_0) t$
 - And a point and normal for each edge
 - P_L, N_L

Cyrus-Beck Algorithm

- Find t such that
- $N_L [P(t) - P_L] = 0$




- Substitute line equation for $P(t)$:
 - $N_L [P_0 + (P_1 - P_0) t - P_L] = 0$
- Solve for t
 - $t = N_L [P_L - P_0] / -N_L [P_1 - P_0]$

Cyrus-Beck Algorithm

- Compute t for line intersection with all four edges
- Discard all ($t < 0$) and ($t > 1$)
- Classify each remaining intersection as
 - Potentially Entering (PE)
 - Potentially Leaving (PL)
- $N_L [P_1 - P_0] > 0$ implies PL
- $N_L [P_1 - P_0] < 0$ implies PE
 - Note that we computed this term when computing t so we can keep it around

Cyrus-Beck Algorithm

- Compute PE with largest t
- Compute PL with smallest t
- Clip to these two points



Cyrus-Beck Algorithm

- Because of horizontal and vertical clip lines:
 - Many computations reduce
- Normals: $(-1, 0)$, $(1, 0)$, $(0, -1)$, $(0, 1)$
- Pick constant points on edges
- solution for t:
 - $-(x_0 - x_{\text{left}}) / (x_1 - x_0)$
 - $(x_0 - x_{\text{right}}) / -(x_1 - x_0)$
 - $-(y_0 - y_{\text{bottom}}) / (y_1 - y_0)$
 - $(y_0 - y_{\text{top}}) / -(y_1 - y_0)$

Comparison


- Cohen-Sutherland
 - Repeated clipping is expensive
 - Best used when trivial acceptance and rejection is possible for most lines
- Cyrus-Beck
 - Computation of t-intersections is cheap
 - Computation of (x,y) clip points is only done once
 - Algorithm doesn't consider trivial accepts/rejects
 - Best when many lines must be clipped
- Liang-Barsky: Optimized Cyrus-Beck
- Nicholl et al.: Fastest, but doesn't do 3D

Clipping Polygons


- Clipping polygons is more complex than clipping the individual lines
 - Input: polygon
 - Output: original polygon, new polygon, or nothing
- The biggest optimizer we had was trivial accept or reject...
- When can we trivially accept/reject a polygon as opposed to the line segments that make up the polygon?

Why Is Clipping Hard?


- What happens to a triangle during clipping?
- Possible outcomes:



triangle \Rightarrow triangle



triangle \Rightarrow quad




triangle \Rightarrow 5-gon

How many sides can a clipped triangle have?


How many sides?

- Seven...




Why Is Clipping Hard?

- A really tough case:



Why Is Clipping Hard?


- A really tough case:



concave polygon \Rightarrow multiple polygons


Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the view window individually
 - Clip the polygon against the view window edge's equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation




Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation



Sutherland-Hodgman Clipping

- Basic idea:
 - Consider each edge of the viewport individually
 - Clip the polygon against the edge equation
 - After doing all edges, the polygon is fully clipped



Sutherland-Hodgman Clipping


- Input/output for algorithm:
 - Input: list of polygon vertices in order
 - Output: list of clipped polygon vertices consisting of old vertices (maybe) and new vertices (maybe)
- Note: this is exactly what we expect from the clipping operation against each edge

Sutherland-Hodgman Clipping

- Sutherland-Hodgman basic routine:
 - Go around polygon one vertex at a time
 - Current vertex has position p
 - Previous vertex had position s , and it has been added to the output if appropriate

Sutherland-Hodgman Clipping

- Edge from s to p takes one of four cases:
(Orange line can be a line or a plane)



Sutherland-Hodgman Clipping

- Four cases:
 - s inside plane and p inside plane
 - Add p to output
 - Note: s has already been added
 - s inside plane and p outside plane
 - Find intersection point i
 - Add i to output
 - s outside plane and p outside plane
 - Add nothing
 - s outside plane and p inside plane
 - Find intersection point i
 - Add i to output, followed by p

Point-to-Plane test

- A very general test to determine if a point p is “inside” a plane P , defined by q and n :




$(p - q) \cdot n < 0$: p inside P

$(p - q) \cdot n = 0$: p on P

$(p - q) \cdot n > 0$: p outside P

Remember: $p \cdot n = |p| |n| \cos(\theta)$

θ = angle between p and n



Finding Line-Plane Intersections

- Edge intersects plane P where $E(t)$ is on P
 - \mathbf{q} is a point on P
 - \mathbf{n} is normal to P

$$(\mathbf{L}(t) - \mathbf{q}) \cdot \mathbf{n} = 0$$

$$(\mathbf{L}_0 + (\mathbf{L}_1 - \mathbf{L}_0) t - \mathbf{q}) \cdot \mathbf{n} = 0$$

$$t = [(\mathbf{q} - \mathbf{L}_0) \cdot \mathbf{n}] / [(\mathbf{L}_1 - \mathbf{L}_0) \cdot \mathbf{n}]$$

- The intersection point $\mathbf{i} = \mathbf{L}(t)$ for this value of t

Projection

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)

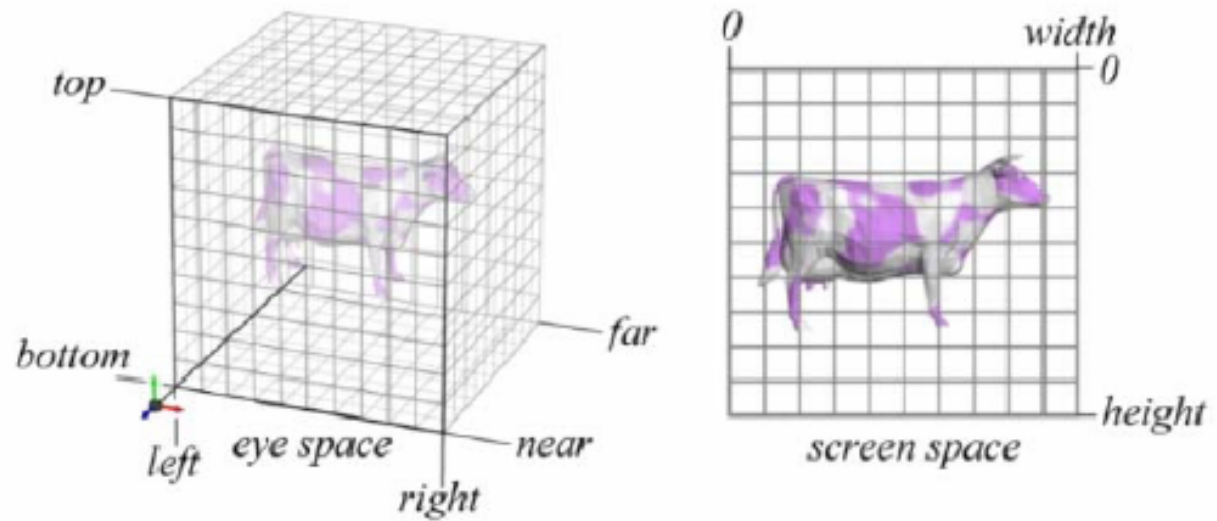
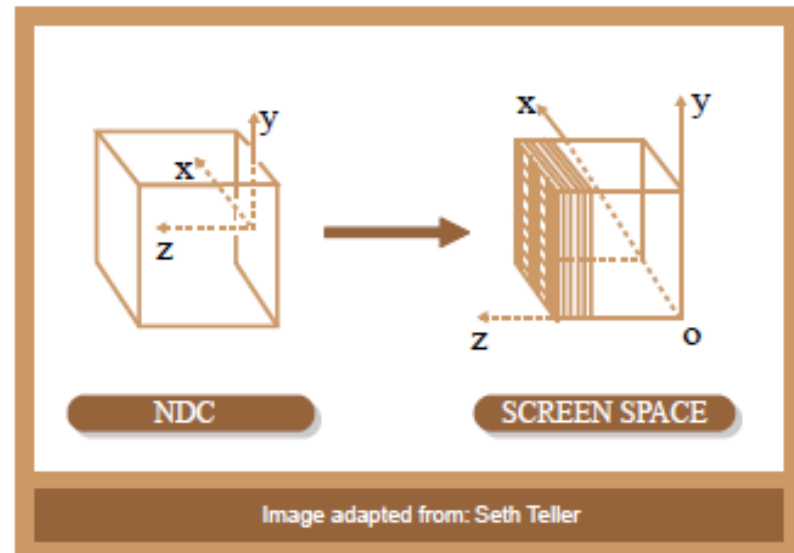
Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- The objects are projected to the 2D image plane (screen space)



Scan Conversion (Rasterization)

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)



Visibility / Display

- Rasterizes objects into pixels
- Interpolate values as we go (color, depth, etc.)




Vector Graphics

- How to generate an image using vectors
 - A line is represented by endpoints (10,10) to (90,90) Cheap transmission
 - The points along the line are computed using a line equation
 - $y = mx + b$ Computation required
 - If you want the image larger, no problem...



Raster Graphics

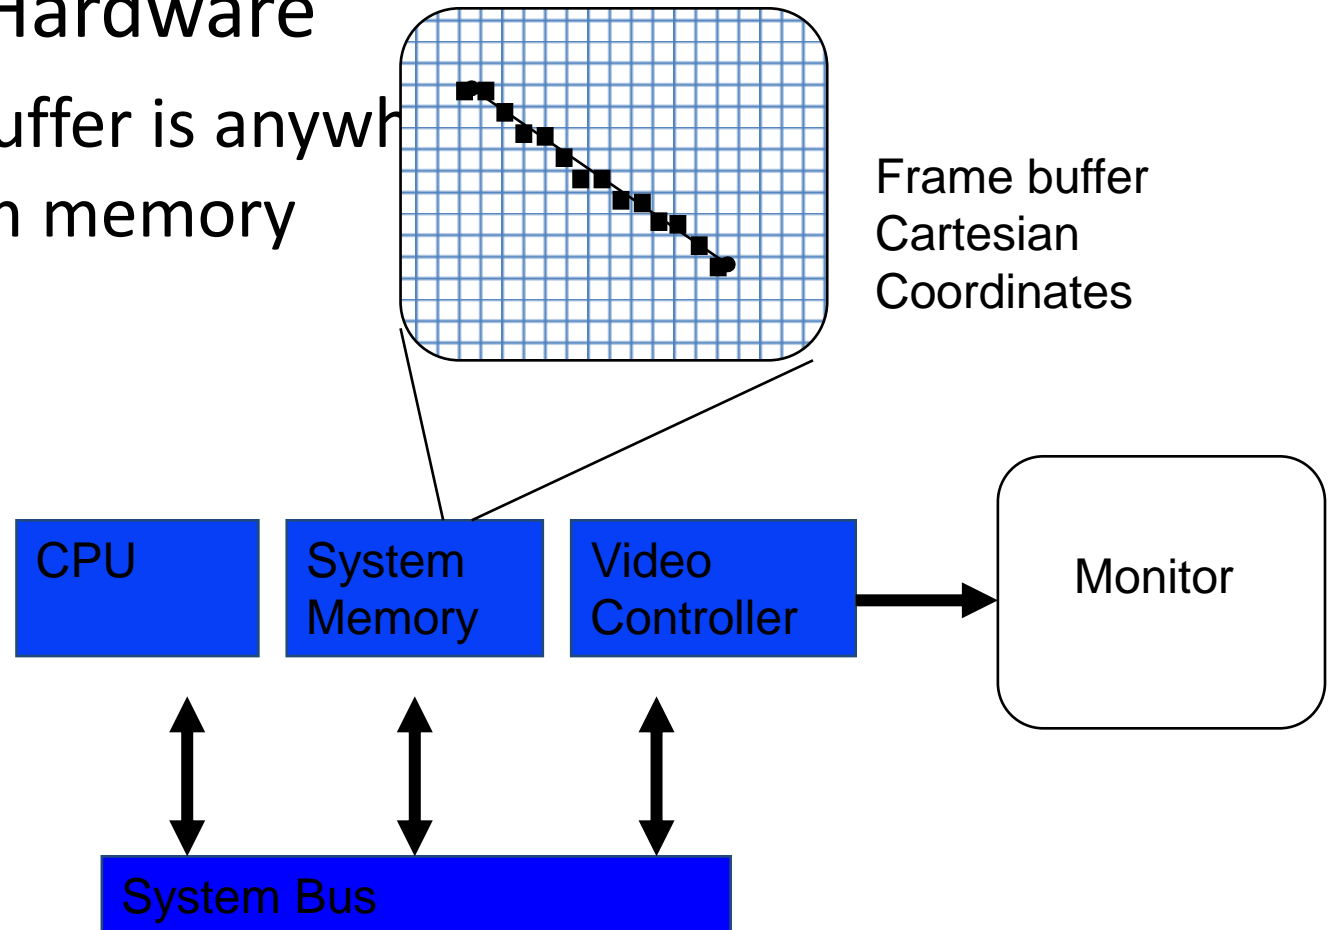
- How to generate a line using rasters
 - A line is represented by assigning some pixels a value of 1
 - Lot's of extra info to communicate
 - The entire line is specified by the pixel values
 - No computation
 - What do we do to make image larger?



Video Controllers

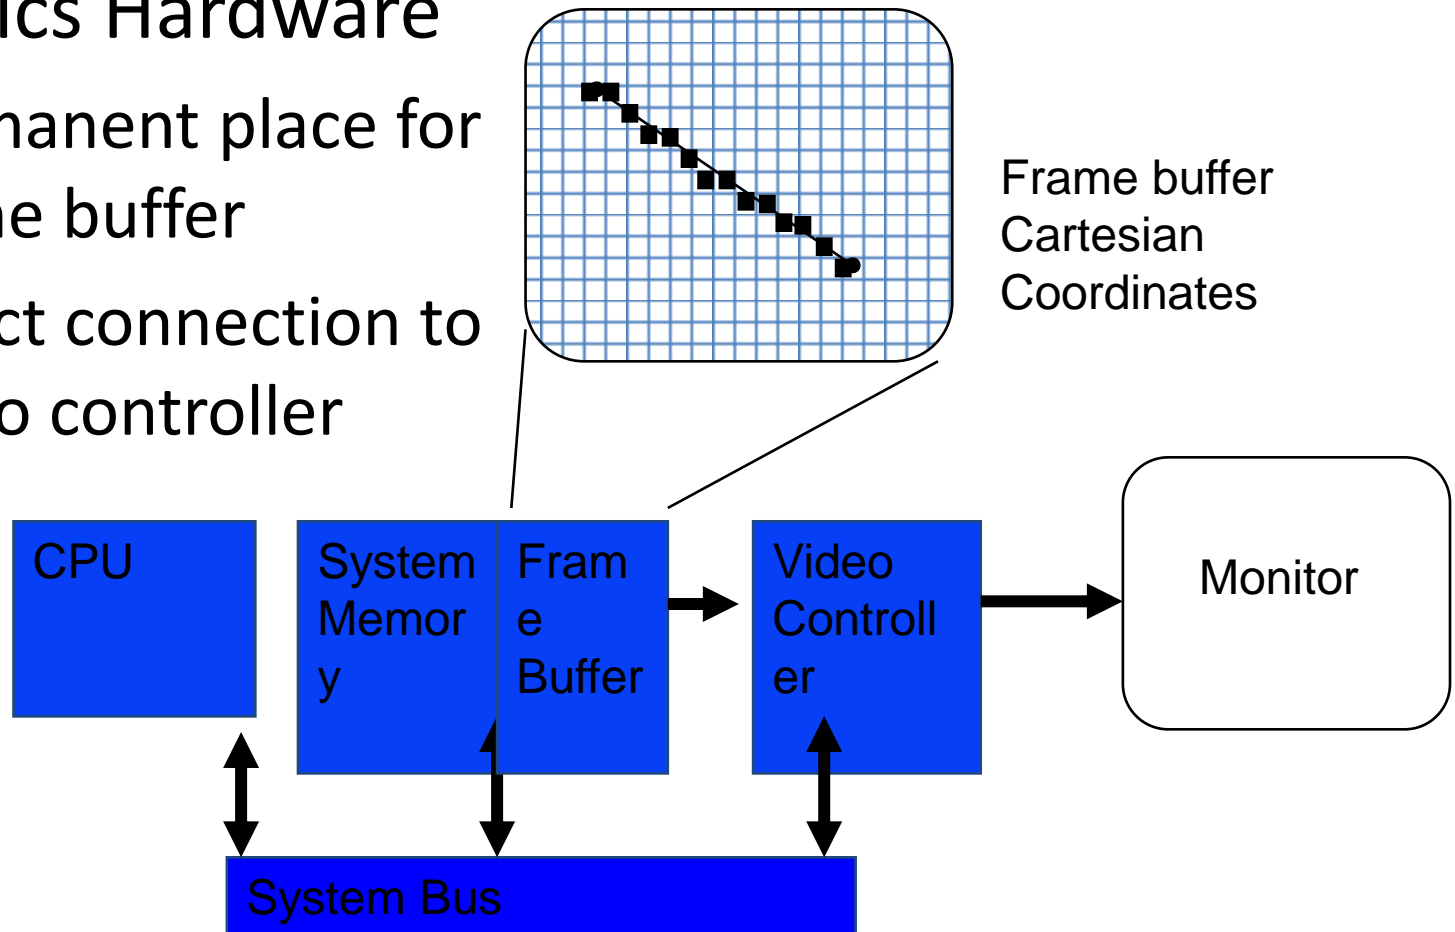
- Graphics Hardware

- Frame buffer is anywhere in system memory



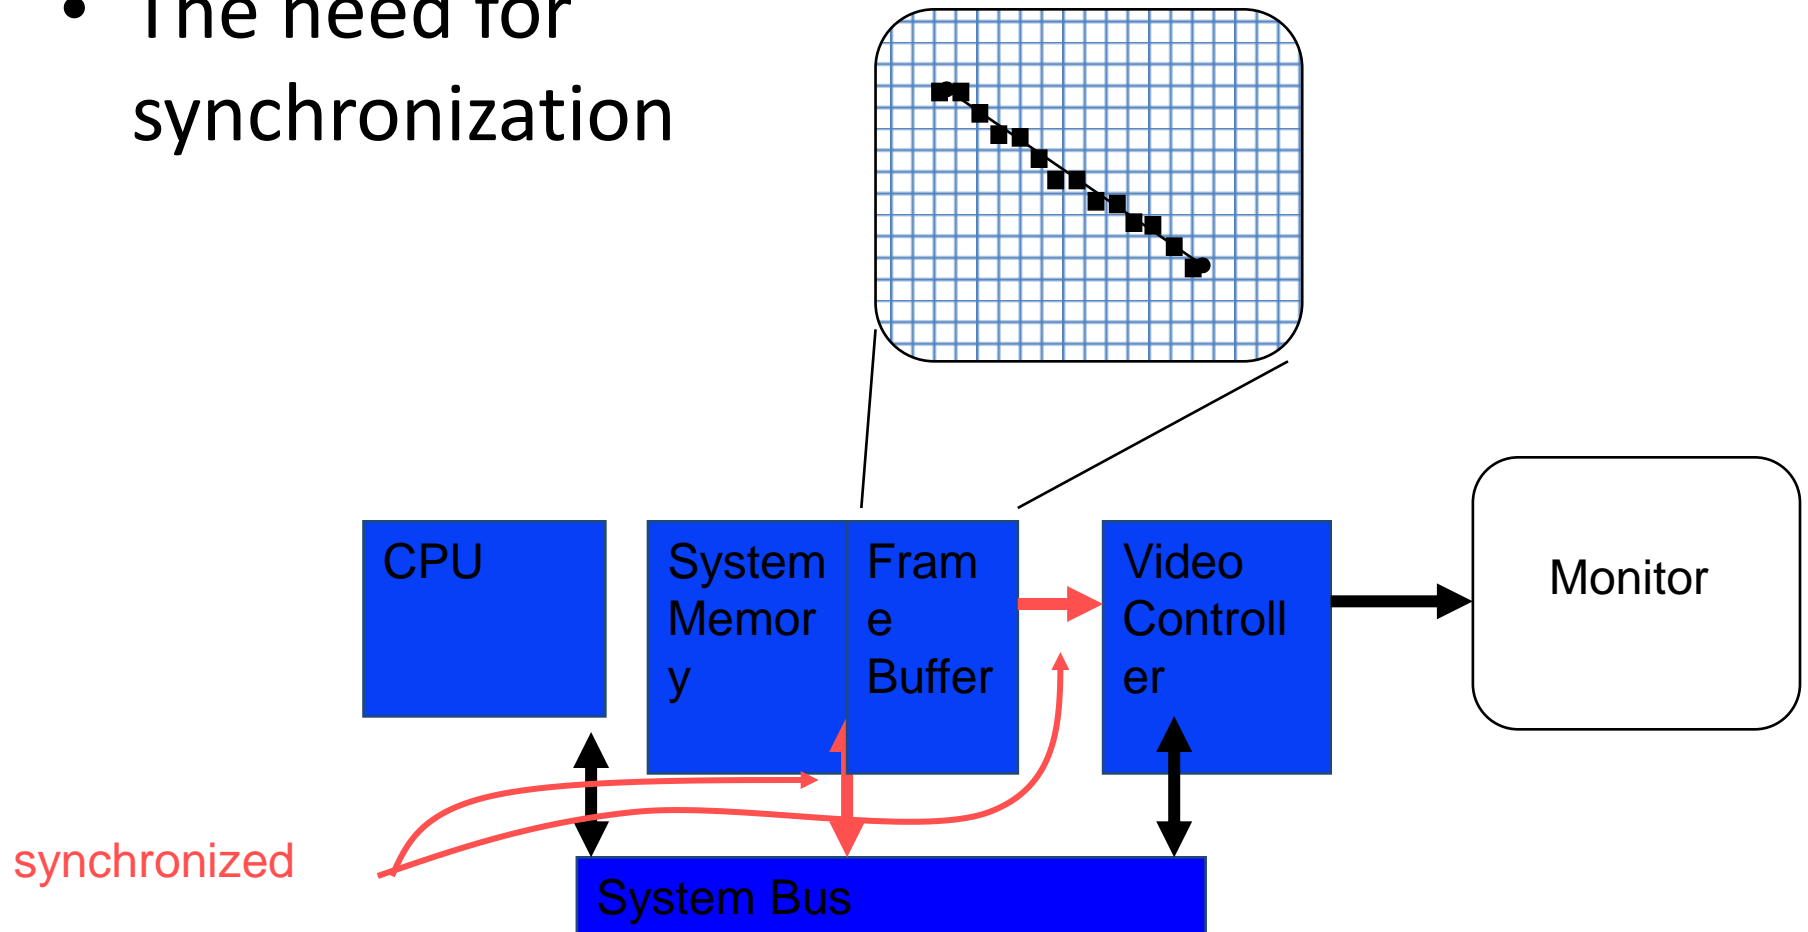
Video Controllers

- Graphics Hardware
 - Permanent place for frame buffer
 - Direct connection to video controller



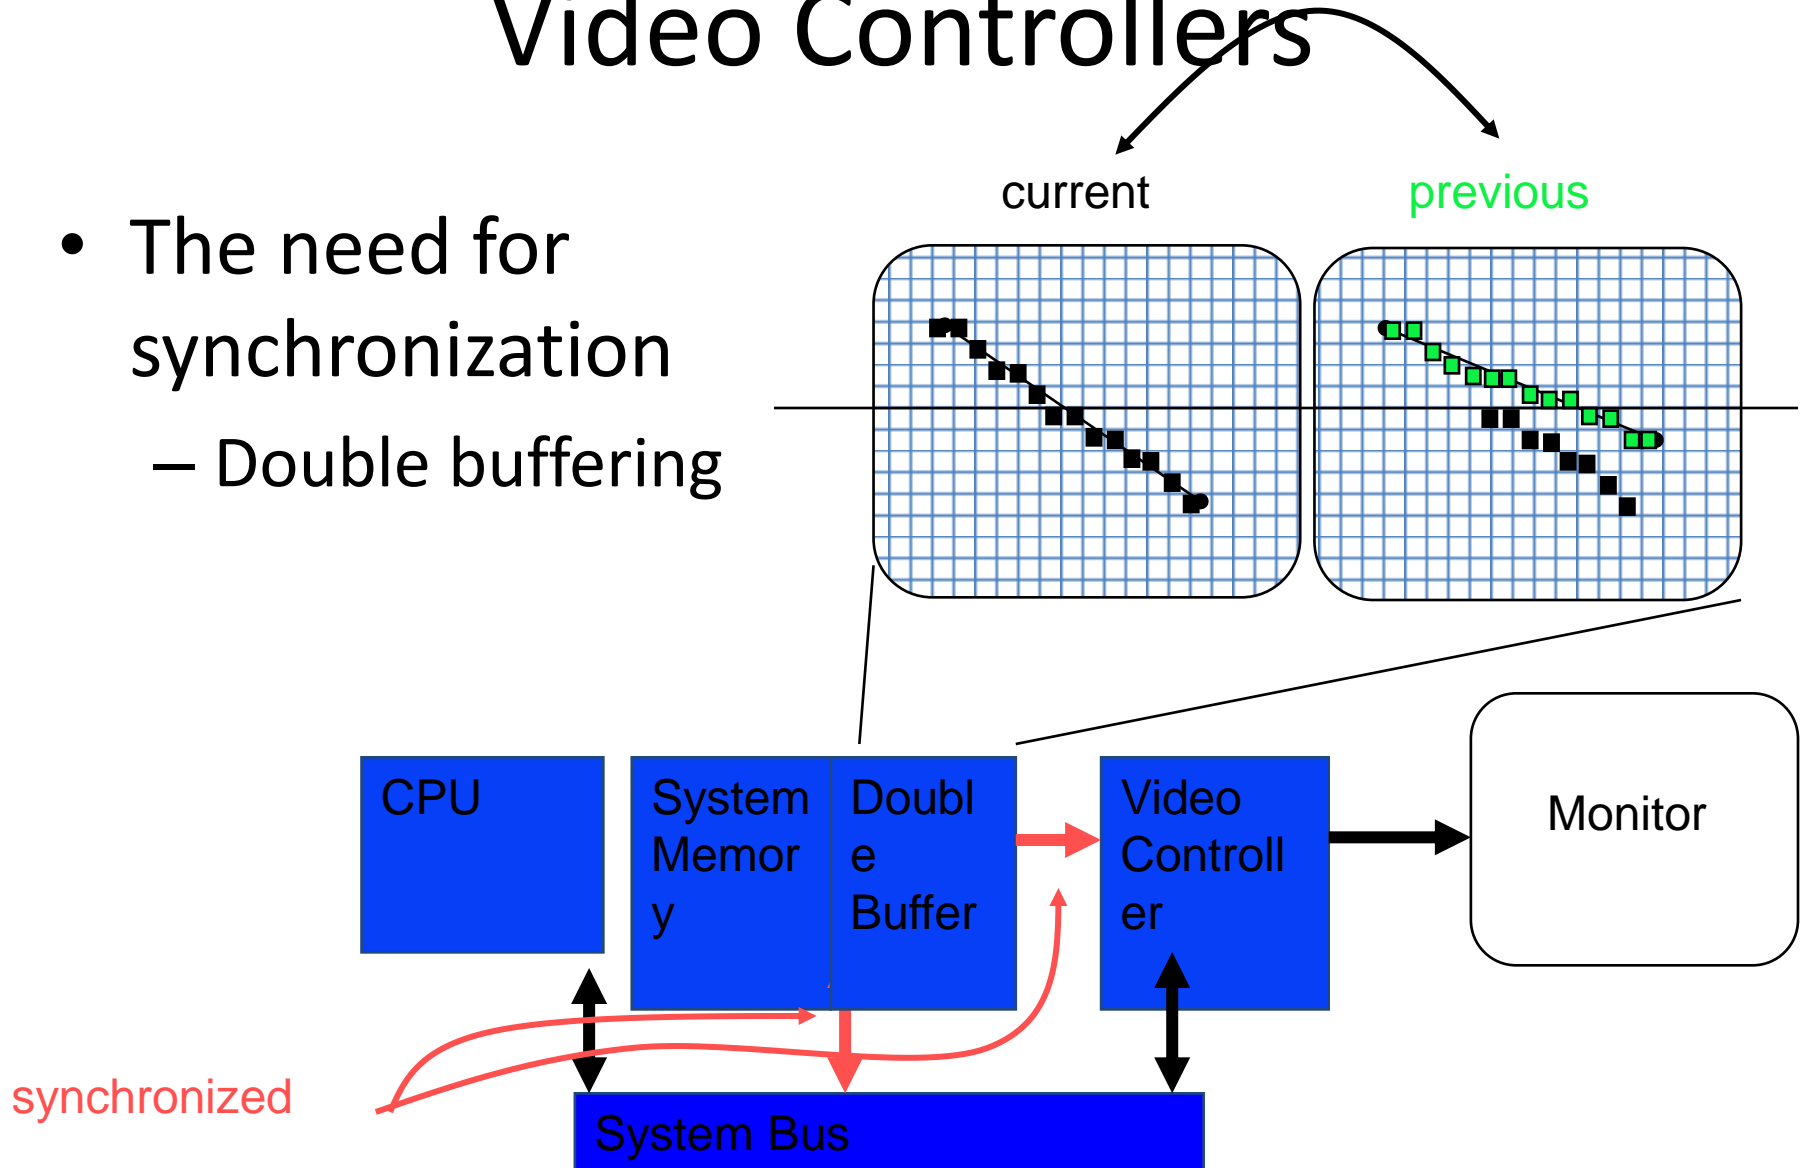
Video Controllers

- The need for synchronization



Video Controllers

- The need for synchronization
 - Double buffering



Frame Buffer

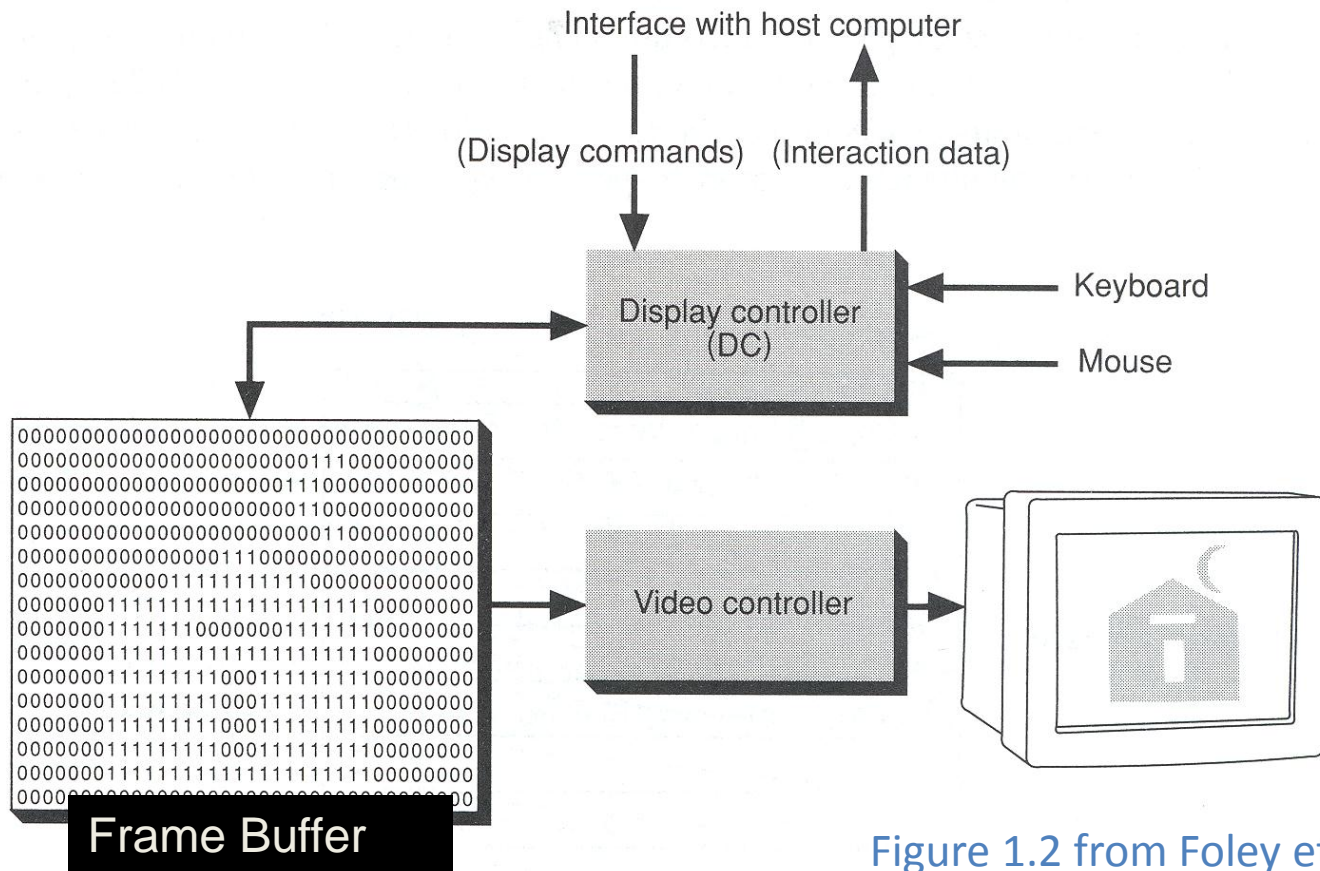



Figure 1.2 from Foley et al.

Frame Buffer Refresh




Refresh rate is usually 30-75Hz

Figure 1.3 from FvDFH


Direct Color Framebuffer

- Store the actual intensities of R, G, and B individually in the framebuffer
- 24 bits per pixel = 8 bits red, 8 bits green, 8 bits blue
 - 16 bits per pixel = ? bits red, ? bits green, ? bits blue



Color Lookup Framebuffer

- Store indices (usually 8 bits) in framebuffer
- Display controller looks up the R,G,B values before triggering the electron guns



Visibility / Display

Modeling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)

Clipping

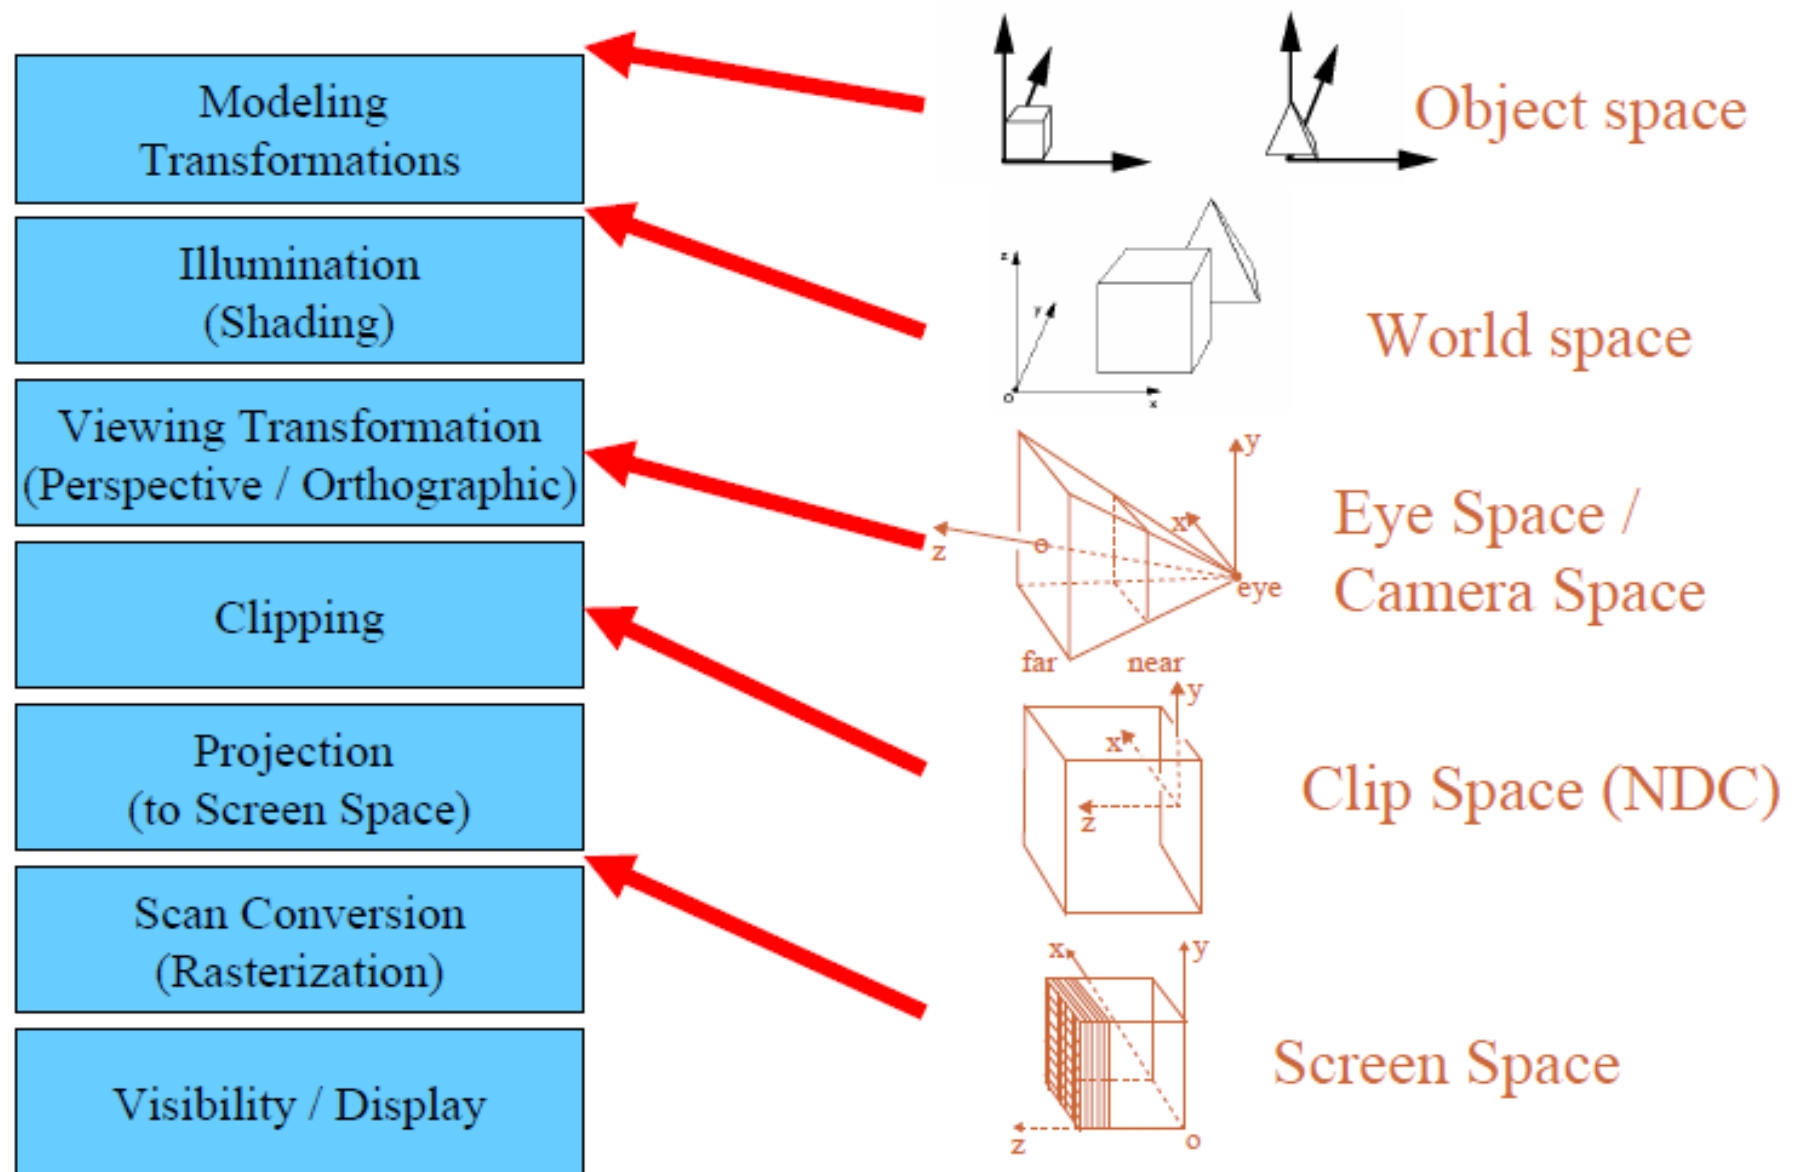
Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Each pixel remembers the closest object (depth buffer)
- Almost every step in the graphics pipeline involves a change of coordinate system. Transformations are central to understanding 3D computer graphics.

Coordinate Systems in the Pipeline



Recap: Rendering Pipeline


- Modeling transformations
- Viewing transformations
- Projection transformations
- Clipping
- Scan conversion
- We now know everything about how to draw a polygon on the screen, except **visible surface determination**

Invisible Primitives

- Why might a polygon be invisible?
 - Polygon outside the *field of view*
 - Polygon is *backfacing*
 - Polygon is *occluded* by object(s) nearer the viewpoint
- For efficiency reasons, we want to avoid spending work on polygons outside field of view or backfacing
- For efficiency and correctness reasons, we need to know when polygons are occluded

View Frustum Clipping

- Remove polygons entirely outside frustum
 - Note that this includes polygons “behind” eye (actually behind near plane)
- Pass through polygons entirely inside frustum
- Modify remaining polygons to include only portions intersecting view frustum




Back-Face Culling

- Most objects in scene are typically “solid”
- More rigorously: **closed, orientable manifolds**
 - Must not cut through itself
 - Must have two distinct sides
 - A sphere is orientable since it has two sides, 'inside' and 'outside'.
 - A Mobius strip or a Klein bottle is not orientable
 - Cannot “walk” from one side to the other
 - A sphere is a closed manifold whereas a plane is not



Back-Face Culling

- On the surface of a closed manifold, polygons whose normals point away from the camera are always occluded:




Back-Face Culling

- Not rendering backfacing polygons improves performance
 - *By how much?*
 - *Reduces by about half the number of polygons to be considered for each pixel*
 - *Every front-facing polygon must have a corresponding rear-facing one*

Occlusion




- For most interesting scenes, some polygons will overlap:



- To render the correct image, we need to determine which polygons **occlude** which

Painter's Algorithm

- Simple approach: render the polygons from back to front, “painting over” previous polygons:




— Draw blue, then green, then orange

- Will this work in the general case?


Painter's Algorithm: Problems

- **Intersecting polygons** present a problem
- Even non-intersecting polygons can form a cycle with no valid visibility order:




Analytic Visibility Algorithms

- Early visibility algorithms computed the set of visible polygon **fragments** directly, then rendered the fragments to a display:



Analytic Visibility Algorithms

- What is the minimum worst-case cost of computing the fragments for a scene composed of n polygons?
- Answer:
 $O(n^2)$
- What's your opinion
- of $O(n^2)$?




Analytic Visibility Algorithms

- So, for about a decade (late 60s to late 70s) there was intense interest in finding efficient algorithms for hidden surface removal
- We'll talk about two:
 - *Binary Space-Partition (BSP) Trees*
 - *Warnock's Algorithm*


Binary Space Partition Trees (1979)

- BSP tree: organize all of space (hence *partition*) into a binary tree
 - *Preprocess*: overlay a binary tree on objects in the scene
 - *Runtime*: correctly traversing this tree enumerates objects from back to front
 - Idea: divide space recursively into half-spaces by choosing *splitting planes*
 - Splitting planes can be arbitrarily oriented


BSP Trees: Objects




BSP Trees: Objects




BSP Trees: Objects



BSP Trees: Objects




BSP Trees: Objects




Rendering BSP Trees

- `renderBSP(BSPtree *T)`
- `BSPtree *near, *far;`
- `if (eye on left side of T->plane)`
- `near = T->left; far = T->right;`
- `else`
- `near = T->right; far = T->left;`
- `renderBSP(far);`
- `if (T is a leaf node)`
- `renderObject(T)`
- `renderBSP(near);`

Rendering BSP Trees



Rendering BSP Trees




Polygons:

BSP Tree Construction

- Split along the plane defined by any polygon from scene
- Classify all polygons into positive or negative half-space of the plane
 - If a polygon intersects plane, split polygon into two and classify them both
- Recurse down the negative half-space
- Recurse down the positive half-space

Discussion: BSP Tree Cons

- No bunnies were harmed in my example
- But what if a splitting plane passes through an object?
 - Split the object; give half to each node




Summary: BSP Trees

- Pros:
 - Simple, elegant scheme
 - Only writes to framebuffer (no reads to see if current polygon is in front of previously rendered polygon, i.e., painters algorithm)
 - Thus very popular for video games (but getting less so)
- Cons:
 - Computationally intense preprocess stage restricts algorithm to static scenes
 - Slow time to construct tree
 - Splitting increases polygon count


Octrees

- Frequently used in modern video games
 - A BSP tree subdivides space into a series of half-spaces using single planes
 - An octree subdivides space into eight voxels using three axis-aligned planes
 - A voxel is labeled as having polygons inside it or not



Octrees

- A voxel may have geometry inside it or subdivide
 - Can have as many as eight children
- Thus we partition 3-D space into 3-D cells
- Checking visibility with polygons now faster due to only checking particular cells
- Quadtrees are a 2-D variant




Warnock's Algorithm (1969)

- Elegant scheme based on a powerful general approach common in graphics: *if the situation is too complex, subdivide*
 - Start with a *root viewport* and a list of all primitives (polygons)
 - Then recursively:
 - Clip objects to viewport
 - If number of objects incident to viewport is zero or one, visibility is trivial
 - Otherwise, subdivide into smaller viewports, distribute primitives among them, and recurse

Warnock's Algorithm

- What is the terminating condition?
- How to determine the correct visible surface in this case?



Warnock's Algorithm


- Pros:
 - Very elegant scheme
 - Extends to any primitive type
- Cons:
 - Hard to embed hierarchical schemes in hardware
 - Complex scenes usually have small polygons and high *depth complexity*
 - Thus most screen regions come down to the single-pixel case

The Z-Buffer Algorithm

- Both BSP trees and Warnock's algorithm were proposed when memory was expensive
 - Example: first 512x512 framebuffer > \$50,000!
- Ed Catmull (mid-70s) proposed a radical new approach called **z-buffering**.
- The big idea: resolve visibility **independently at each pixel**



The Z-Buffer Algorithm

- We know how to rasterize polygons into an image discretized into pixels:



The Z-Buffer Algorithm

- What happens if multiple primitives occupy the same pixel on the screen? Which is allowed to paint the pixel?



The Z-Buffer Algorithm

- Idea: retain depth (Z in eye coordinates) through projection transform
 - Use canonical viewing volumes
 - Each vertex has z coordinate (relative to eye point) intact

The Z-Buffer Algorithm

- Augment framebuffer with **Z-buffer** or **depth buffer** which stores Z value at each pixel
 - At frame beginning, initialize all pixel depths to ∞
 - When rasterizing, interpolate depth (Z) across polygon and store in pixel of Z-buffer
 - Suppress writing to a pixel if its Z value is more distant than the Z value already stored there

The Z-Buffer Algorithm

- How much memory does the Z-buffer use?
- Does the image rendered depend on the drawing order?
- Does the time to render the image depend on the drawing order?
- How does Z-buffer load scale with visible polygons? With framebuffer resolution?

Z-Buffer Pros

- Simple!!!
- Easy to implement in hardware
- Polygons can be processed in arbitrary order
- Easily handles polygon interpenetration
- Enables deferred shading
 - Rasterize shading parameters (e.g., surface normal) and only shade final visible fragments

Z-Buffer Cons

- Lots of memory (e.g. 1280x1024x32 bits)
 - With 16 bits cannot discern millimeter differences in objects at 1 km distance
- Read-Modify-Write in inner loop requires fast memory
- Hard to do analytic antialiasing
 - We don't know which polygon to map pixel back to
- Shared edges are handled inconsistently
 - ***Ordering dependent***
- Hard to simulate translucent polygons
 - We throw away color of polygons behind closest one