

Introduction to Information Retrieval

<http://informationretrieval.org>

IIR 3: Dictionaries and tolerant retrieval

Hinrich Schütze

Institute for Natural Language Processing, Universität Stuttgart

2008.04.29

Overview

- 1 Recap
- 2 Dictionaries
- 3 Wildcard queries
- 4 Spelling correction
- 5 Soundex

Outline

- 1 Recap
- 2 Dictionaries
- 3 Wildcard queries
- 4 Spelling correction
- 5 Soundex

Type/token distinction

- **Token** – An instance of a word or term occurring in a document.

Type/token distinction

- **Token** – An instance of a word or term occurring in a document.
- **Type** – An equivalence class of tokens.

Type/token distinction

- **Token** – An instance of a word or term occurring in a document.
- **Type** – An equivalence class of tokens.
- *In June, the dog likes to chase the cat in the barn.*

Type/token distinction

- **Token** – An instance of a word or term occurring in a document.
- **Type** – An equivalence class of tokens.
- *In June, the dog likes to chase the cat in the barn.*
- How many tokens? How many types?

Type/token distinction

- **Token** – An instance of a word or term occurring in a document.
- **Type** – An equivalence class of tokens.
- *In June, the dog likes to chase the cat in the barn.*
- How many tokens? How many types?
- 12 tokens, 9 types

Problems in tokenization

- What are the delimiters? Space? Apostrophe? Hyphen?

Problems in tokenization

- What are the delimiters? Space? Apostrophe? Hyphen?
- For each of these: sometimes they delimit, sometimes they don't.

Problems in tokenization

- What are the delimiters? Space? Apostrophe? Hyphen?
- For each of these: sometimes they delimit, sometimes they don't.
- No whitespace in many languages! (e.g., Chinese)

Problems in tokenization

- What are the delimiters? Space? Apostrophe? Hyphen?
- For each of these: sometimes they delimit, sometimes they don't.
- No whitespace in many languages! (e.g., Chinese)
- No whitespace in Dutch, German, Swedish compounds (*Lebensversicherungsgesellschaftsangestellter*)

Problems in tokenization

- What are the delimiters? Space? Apostrophe? Hyphen?
- For each of these: sometimes they delimit, sometimes they don't.
- No whitespace in many languages! (e.g., Chinese)
- No whitespace in Dutch, German, Swedish compounds (*Lebensversicherungsgesellschaftsangestellter*)
- No whitespace in English: *database*, *whitespace*

Problems in “equivalence classing”

- A term is an equivalence class of tokens.

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding
- Stemming, Porter stemmer

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding
- Stemming, Porter stemmer
- Morphological analysis: inflectional vs. derivational

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding
- Stemming, Porter stemmer
- Morphological analysis: inflectional vs. derivational
- Equivalence classing problems in other languages

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding
- Stemming, Porter stemmer
- Morphological analysis: inflectional vs. derivational
- Equivalence classing problems in other languages
 - More complex morphology than in English

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding
- Stemming, Porter stemmer
- Morphological analysis: inflectional vs. derivational
- Equivalence classing problems in other languages
 - More complex morphology than in English
 - Finnish: a single verb may have 12,000 different forms


Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding
- Stemming, Porter stemmer
- Morphological analysis: inflectional vs. derivational
- Equivalence classing problems in other languages
 - More complex morphology than in English
 - Finnish: a single verb may have 12,000 different forms
 - Words written in different alphabets (Hiragana vs. Chinese characters)

Problems in “equivalence classing”

- A term is an equivalence class of tokens.
- How do we define equivalence classes?
- Numbers (3/12/91 vs. 12/3/91)
- Case folding
- Stemming, Porter stemmer
- Morphological analysis: inflectional vs. derivational
- Equivalence classing problems in other languages
 - More complex morphology than in English
 - Finnish: a single verb may have 12,000 different forms
 - Words written in different alphabets (Hiragana vs. Chinese characters)
 - Accents, umlauts

Skip pointers



Positional indexes

- Postings lists in a **positional index**: each posting is a docID and a **list of positions**
- Example: to_1 be_2 or_3 not_4 to_5 be_6

TO, 993427:

$\langle 1, 6: \langle 7, 18, 33, 72, 86, 231 \rangle;$
 $2, 5: \langle 1, 17, 74, 222, 255 \rangle;$
 $4, 5: \langle 8, 16, 190, 429, 433 \rangle;$
 $5, 2: \langle 363, 367 \rangle;$
 $7, 3: \langle 13, 23, 191 \rangle; \dots \rangle$

BE, 178239:

$\langle 1, 2: \langle 17, 25 \rangle;$
 $4, 5: \langle 17, 191, 291, 430, 434 \rangle;$
 $5, 3: \langle 14, 19, 101 \rangle; \dots \rangle$

Document 4 is a match.

Positional indexes

- Postings lists in a **positional index**: each posting is a docID and a **list of positions**
- **Example:** *to₁ be₂ or₃ not₄ to₅ be₆*

TO, 993427:

⟨ 1, 6: ⟨7, 18, 33, 72, 86, 231⟩;
 2, 5: ⟨1, 17, 74, 222, 255⟩;
 4, 5: ⟨8, 16, 190, **429**, **433**⟩;
 5, 2: ⟨363, 367⟩;
 7, 3: ⟨13, 23, 191⟩; ... ⟩

BE, 178239:

⟨ 1, 2: ⟨17, 25⟩;
 4, 5: ⟨17, 191, 291, **430**, **434**⟩;
 5, 3: ⟨14, 19, 101⟩; ... ⟩

Document 4 is a match.

Positional indexes

- With a positional index, we can answer **phrase queries**.

Positional indexes

- With a positional index, we can answer **phrase queries**.
- With a positional index, we can answer **proximity queries**.

Outline

- 1 Recap
- 2 Dictionaries
- 3 Wildcard queries
- 4 Spelling correction
- 5 Soundex

Inverted index

For each term t , we store a list of all documents that contain t .

BRUTUS	→	1	2	4	11	31	45	173	174
--------	---	---	---	---	----	----	----	-----	-----

CAESAR	→	1	2	4	5	6	16	57	132	...
--------	---	---	---	---	---	---	----	----	-----	-----

CALPURNIA	→	2	31	54	101
-----------	---	---	----	----	-----

⋮

dictionary

postings

Inverted index

For each term t , we store a list of all documents that contain t .

BRUTUS	→	1	2	4	11	31	45	173	174
--------	---	---	---	---	----	----	----	-----	-----

CAESAR	→	1	2	4	5	6	16	57	132	...
--------	---	---	---	---	---	---	----	----	-----	-----

CALPURNIA	→	2	31	54	101
-----------	---	---	----	----	-----

⋮

dictionary

postings

Dictionaries

- The dictionary is the data structure for storing the term vocabulary.

Dictionaries

- The dictionary is the data structure for storing the term vocabulary.
- Term vocabulary: the data

Dictionaries

- The dictionary is the data structure for storing the term vocabulary.
- Term vocabulary: the data
- Dictionary: the data structure for storing the term vocabulary

Dictionary as array of fixed-width entries

- For each term, we need to store a couple of items:

Dictionary as array of fixed-width entries

- For each term, we need to store a couple of items:
 - document frequency

Dictionary as array of fixed-width entries

- For each term, we need to store a couple of items:
 - document frequency
 - pointer to postings list

Dictionary as array of fixed-width entries

- For each term, we need to store a couple of items:
 - document frequency
 - pointer to postings list
 - ...

Dictionary as array of fixed-width entries

- For each term, we need to store a couple of items:
 - document frequency
 - pointer to postings list
 - ...
- Assume for the time being that we can store this information in a fixed-length entry.

Dictionary as array of fixed-width entries

- For each term, we need to store a couple of items:
 - document frequency
 - pointer to postings list
 - ...
- Assume for the time being that we can store this information in a fixed-length entry.
- Assume that we store these entries in an array.

Dictionary as array of fixed-width entries

term	document frequency	pointer to postings list
a	656,265	→
aachen	65	→
...
zulu	221	→

space needed: 20 bytes 4 bytes 4 bytes

How do we look up an element in this array at query time?

Data structures for looking up term

- Two main classes of data structures: hashes and trees

Data structures for looking up term

- Two main classes of data structures: hashes and trees
- Some IR systems use hashes, some use trees.

Data structures for looking up term

- Two main classes of data structures: hashes and trees
- Some IR systems use hashes, some use trees.
- Criteria for when to use hashes vs. trees:

Data structures for looking up term

- Two main classes of data structures: hashes and trees
- Some IR systems use hashes, some use trees.
- Criteria for when to use hashes vs. trees:
 - Is there a fixed number of terms or will it keep growing?

Data structures for looking up term

- Two main classes of data structures: hashes and trees
- Some IR systems use hashes, some use trees.
- Criteria for when to use hashes vs. trees:
 - Is there a fixed number of terms or will it keep growing?
 - What are the relative frequencies with which various keys will be accessed?

Data structures for looking up term

- Two main classes of data structures: hashes and trees
- Some IR systems use hashes, some use trees.
- Criteria for when to use hashes vs. trees:
 - Is there a fixed number of terms or will it keep growing?
 - What are the relative frequencies with which various keys will be accessed?
 - How many terms are we likely to have?

Hashes

- Each vocabulary term is hashed into an integer.

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions
- At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions
- At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array
- Pros: Lookup in a hash is faster than lookup in a tree.

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions
- At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array
- Pros: Lookup in a hash is faster than lookup in a tree.
- Cons

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions
- At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array
- Pros: Lookup in a hash is faster than lookup in a tree.
- Cons
 - no way to find minor variants (*resume* vs. *résumé*)

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions
- At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array
- Pros: Lookup in a hash is faster than lookup in a tree.
- Cons
 - no way to find minor variants (*resume* vs. *résumé*)
 - no prefix search (all terms starting with *automat*)

Hashes

- Each vocabulary term is hashed into an integer.
- Try to avoid collisions
- At query time, do the following: hash query term, resolve collisions, locate entry in fixed-width array
- Pros: Lookup in a hash is faster than lookup in a tree.
- Cons
 - no way to find minor variants (*resume* vs. *résumé*)
 - no prefix search (all terms starting with *automat*)
 - need to rehash everything periodically if vocabulary keeps growing

Trees

- Trees solve the prefix problem (find all terms starting with *automat*).

Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree

Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree
- Search is slightly slower than in hashes: $O(\log M)$, where M is the size of the vocabulary.

Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree
- Search is slightly slower than in hashes: $O(\log M)$, where M is the size of the vocabulary.
- $O(\log M)$ only holds for **balanced** trees.

Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree
- Search is slightly slower than in hashes: $O(\log M)$, where M is the size of the vocabulary.
- $O(\log M)$ only holds for **balanced** trees.
- Rebalancing binary trees is expensive.

Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree
- Search is slightly slower than in hashes: $O(\log M)$, where M is the size of the vocabulary.
- $O(\log M)$ only holds for **balanced** trees.
- Rebalancing binary trees is expensive.
- **B-trees** mitigate the rebalancing problem.


Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree
- Search is slightly slower than in hashes: $O(\log M)$, where M is the size of the vocabulary.
- $O(\log M)$ only holds for **balanced** trees.
- Rebalancing binary trees is expensive.
- **B-trees** mitigate the rebalancing problem.
- B-tree definition: every internal node has a number of children in the interval $[a, b]$ where a, b are appropriate positive integers, e.g., $[2, 4]$.


Trees

- Trees solve the prefix problem (find all terms starting with *automat*).
- Simplest tree: binary tree
- Search is slightly slower than in hashes: $O(\log M)$, where M is the size of the vocabulary.
- $O(\log M)$ only holds for **balanced** trees.
- Rebalancing binary trees is expensive.
- **B-trees** mitigate the rebalancing problem.
- B-tree definition: every internal node has a number of children in the interval $[a, b]$ where a, b are appropriate positive integers, e.g., $[2, 4]$.
- Note that we need a standard ordering for characters in order to be able to use trees.

Binary tree



B-tree



Outline

- 1 Recap
- 2 Dictionaries
- 3 Wildcard queries
- 4 Spelling correction
- 5 Soundex

Wildcard queries

- `mon*`: find all docs containing any term beginning with *mon*

Wildcard queries

- mon^* : find all docs containing any term beginning with *mon*
- Easy with B-tree dictionary: retrieve all terms t in the range:
 $\text{mon} \leq t < \text{moo}$

Wildcard queries

- mon^* : find all docs containing any term beginning with *mon*
- Easy with B-tree dictionary: retrieve all terms t in the range:
 $mon \leq t < moo$
- $*mon$: find all docs containing any term ending with *mon*

Wildcard queries

- mon^* : find all docs containing any term beginning with *mon*
- Easy with B-tree dictionary: retrieve all terms t in the range:
 $mon \leq t < moo$
- $*mon$: find all docs containing any term ending with *mon*
 - Maintain an additional tree for terms *backwards*

Wildcard queries

- mon^* : find all docs containing any term beginning with *mon*
- Easy with B-tree dictionary: retrieve all terms t in the range:
 $mon \leq t < moo$
- $*mon$: find all docs containing any term ending with *mon*
 - Maintain an additional tree for terms *backwards*
 - Then retrieve all terms t in the range: $nom \leq t < non$

Query processing

- At this point, we have an enumeration of all terms in the dictionary that match the wildcard query.

Query processing

- At this point, we have an enumeration of all terms in the dictionary that match the wildcard query.
- We still have to look up the postings for each enumerated term.

Query processing

- At this point, we have an enumeration of all terms in the dictionary that match the wildcard query.
- We still have to look up the postings for each enumerated term.
- E.g., consider the query: `gen* AND universit*`

Query processing

- At this point, we have an enumeration of all terms in the dictionary that match the wildcard query.
- We still have to look up the postings for each enumerated term.
- E.g., consider the query: `gen* AND universit*`
- This may result in the execution of many Boolean `AND` queries.

How to handle * in the middle of a term

- Example: m*nchen

How to handle * in the middle of a term

- Example: m*nchen
- We could look up m* and *nchen in the B-tree and intersect the two term sets.

How to handle * in the middle of a term

- Example: m*nchen
- We could look up m* and *nchen in the B-tree and intersect the two term sets.
- Expensive

How to handle * in the middle of a term

- Example: m*nchen
- We could look up m* and *nchen in the B-tree and intersect the two term sets.
- Expensive
- Alternative: [permuterm](#) index

How to handle * in the middle of a term

- Example: m*nchen
- We could look up m* and *nchen in the B-tree and intersect the two term sets.
- Expensive
- Alternative: [permuterm](#) index
- Basic idea: Rotate every wildcard query, so that the * occurs at the end.


Permuterm index

- For term HELLO: add *hello\$*, *ello\$h*, *llo\$he*, *lo\$hel*, and *o\$hell* to the B-tree where \$ is a special symbol

Permuterm index

- For term HELLO: add *hello\$*, *ello\$h*, *llo\$he*, *lo\$hel*, and *o\$hell* to the B-tree where \$ is a special symbol
- Queries

Permuterm \rightarrow term mapping



Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries

Permuterm index

- For HELLO, we've stored: *hello\$*, *ello\$h*, *llo\$he*, *lo\$hel*, and *o\$hell*
- Queries
 - For X, look up X\$

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$
 - For *X, look up X\$*

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$
 - For *X, look up X\$*
 - For *X*, look up X*

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$
 - For *X, look up X\$*
 - For *X*, look up X*
 - For X*Y, look up Y\$X*

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$
 - For *X, look up X\$*
 - For *X*, look up X*
 - For X*Y, look up Y\$X*
 - Example: For hel*o, look up o\$hel*

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$
 - For *X, look up X\$*
 - For *X*, look up X*
 - For X*Y, look up Y\$X*
 - Example: For hel*o, look up o\$hel*
 - How do we handle X*Y*Z?

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$
 - For *X, look up X\$*
 - For *X*, look up X*
 - For X*Y, look up Y\$X*
 - Example: For hel*o, look up o\$hel*
 - How do we handle X*Y*Z?
- It's really a tree and should be called permuterm tree.

Permuterm index

- For HELLO, we've stored: *hello\$, ello\$h, llo\$he, lo\$hel, and o\$hell*
- Queries
 - For X, look up X\$
 - For X*, look up X*\$
 - For *X, look up X\$*
 - For *X*, look up X*
 - For X*Y, look up Y\$X*
 - Example: For hel*o, look up o\$hel*
 - How do we handle X*Y*Z?
- It's really a tree and should be called permuterm tree.
- But permuterm index is more common name.

Processing a lookup in the permuterm index

- Rotate query wildcard to the right

Processing a lookup in the permuterm index

- Rotate query wildcard to the right
- Use B-tree lookup as before

Processing a lookup in the permuterm index

- Rotate query wildcard to the right
- Use B-tree lookup as before
- Problem: Permuterm **quadruples** the size of the dictionary compared to a regular B-tree. (empirical number)

k -gram indexes

- More space-efficient than permuterm index

k -gram indexes

- More space-efficient than permuterm index
- Enumerate all character k -grams (sequence of k characters) occurring in a term

k -gram indexes

- More space-efficient than permuterm index
- Enumerate all character k -grams (sequence of k characters) occurring in a term
- 2-grams are called **bigrams**.

k -gram indexes

- More space-efficient than permuterm index
- Enumerate all character k -grams (sequence of k characters) occurring in a term
- 2-grams are called **bigrams**.
- Example: from *April is the cruelest month* we get the bigrams:
 $\$a\ ap\ pr\ ri\ il\ l\$ \$i\ is\ s\$ \$t\ th\ he\ e\$ \$c\ cr\ ru\ ue\ el\ le\ es\ st\ t\$ \m
 $mo\ on\ nt\ h\$$

k -gram indexes

- More space-efficient than permuterm index
- Enumerate all character k -grams (sequence of k characters) occurring in a term
- 2-grams are called **bigrams**.
- Example: from *April is the cruelest month* we get the bigrams:
 $\$a\ ap\ pr\ ri\ il\ l\$ \$i\ is\ s\$ \$t\ th\ he\ e\$ \$c\ cr\ ru\ ue\ el\ le\ es\ st\ t\$ \m
 $mo\ on\ nt\ h\$$
- $\$$ is a special word boundary symbol.

k -gram indexes

- More space-efficient than permuterm index
- Enumerate all character k -grams (sequence of k characters) occurring in a term
- 2-grams are called **bigrams**.
- Example: from *April is the cruelest month* we get the bigrams:
 $\$a\ ap\ pr\ ri\ il\ l\$ \$i\ is\ s\$ \$t\ th\ he\ e\$ \$c\ cr\ ru\ ue\ el\ le\ es\ st\ t\$ \m
 $mo\ on\ nt\ h\$$
- $\$$ is a special word boundary symbol.
- Maintain an inverted index from bigrams to the terms that contain the bigram

Postings list in a 3-gram index



Bigram indexes

- Note that we now have two different types of inverted indexes

Bigram indexes

- Note that we now have two different types of inverted indexes
- The term-document inverted index for finding documents based on a query consisting of terms

Bigram indexes

- Note that we now have two different types of inverted indexes
- The term-document inverted index for finding documents based on a query consisting of terms
- The k -gram index for finding terms based on a query consisting of k -grams

Processing wildcarded terms in a bigram index

- Query mon* can now be run as:
\$m AND mo AND on

Processing wildcarded terms in a bigram index

- Query `mon*` can now be run as:
`$m AND mo AND on`
- Gets us all terms with the prefix *mon* . . .

Processing wildcarded terms in a bigram index

- Query `mon*` can now be run as:
`$m AND mo AND on`
- Gets us all terms with the prefix *mon* ...
- ... but also many “false positives” like `MOON`.

Processing wildcard terms in a bigram index

- Query `mon*` can now be run as:
 `$m AND mo AND on`
- Gets us all terms with the prefix *mon* . . .
- . . . but also many “false positives” like `MOON`.
- We must postfilter these terms against query.

Processing wildcarded terms in a bigram index

- Query `mon*` can now be run as:
`$m AND mo AND on`
- Gets us all terms with the prefix *mon* . . .
- . . . but also many “false positives” like `MOON`.
- We must postfilter these terms against query.
- Surviving terms are then looked up in the term-document inverted index.

Processing wildcarded terms in a bigram index

- Query `mon*` can now be run as:
`$m AND mo AND on`
- Gets us all terms with the prefix *mon* . . .
- . . . but also many “false positives” like `MOON`.
- We must postfilter these terms against query.
- Surviving terms are then looked up in the term-document inverted index.
- *k*-gram indexes are fast and space efficient (compared to permuterm indexes).

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`
- Most straightforward semantics: Conjunction of disjunctions

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`
- Most straightforward semantics: Conjunction of disjunctions
- Very expensive

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`
- Most straightforward semantics: Conjunction of disjunctions
- Very expensive
- Does Google allow wildcard queries?

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`
- Most straightforward semantics: Conjunction of disjunctions
- Very expensive
- Does Google allow wildcard queries?
- Why?

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`
- Most straightforward semantics: Conjunction of disjunctions
- Very expensive
- Does Google allow wildcard queries?
- Why?
- Users hate to type.

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`
- Most straightforward semantics: Conjunction of disjunctions
- Very expensive
- Does Google allow wildcard queries?
- Why?
- Users hate to type.
- If abbreviated queries like `pyth* theo*` for pythagoras' theorem are legal, users will use them ...

Processing wildcard queries in the term-document index

- As before, we must potentially execute a large number of Boolean queries for each enumerated, filtered term.
- Recall the query: `gen* AND universit*`
- Most straightforward semantics: Conjunction of disjunctions
- Very expensive
- Does Google allow wildcard queries?
- Why?
- Users hate to type.
- If abbreviated queries like `pyth* theo*` for pythagoras' theorem are legal, users will use them ...
- ... a lot

Outline

- 1 Recap
- 2 Dictionaries
- 3 Wildcard queries
- 4 Spelling correction**
- 5 Soundex

Spelling correction

- Two principal uses

Spelling correction

- Two principal uses
 - Correcting documents being indexed

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries
- Two different methods for spelling correction

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries
- Two different methods for spelling correction
- **Isolated word** spelling correction

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries
- Two different methods for spelling correction
- **Isolated word** spelling correction
 - Check each word on its own for misspelling

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries
- Two different methods for spelling correction
- **Isolated word** spelling correction
 - Check each word on its own for misspelling
 - Will not catch typos resulting in correctly spelled words, e.g.,
*an asteroid that fell **form** the sky*

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries
- Two different methods for spelling correction
- **Isolated word** spelling correction
 - Check each word on its own for misspelling
 - Will not catch typos resulting in correctly spelled words, e.g.,
*an asteroid that fell **form** the sky*
- **Context-sensitive** spelling correction

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries
- Two different methods for spelling correction
- **Isolated word** spelling correction
 - Check each word on its own for misspelling
 - Will not catch typos resulting in correctly spelled words, e.g.,
*an asteroid that fell **form** the sky*
- **Context-sensitive** spelling correction
 - Look at surrounding words

Spelling correction

- Two principal uses
 - Correcting documents being indexed
 - Correcting user queries
- Two different methods for spelling correction
- **Isolated word** spelling correction
 - Check each word on its own for misspelling
 - Will not catch typos resulting in correctly spelled words, e.g.,
*an asteroid that fell **form** the sky*
- **Context-sensitive** spelling correction
 - Look at surrounding words
 - Can correct *form/from* error above

Correcting documents

- We're not interested in interactive spelling correction of documents (e.g., MS Word) in this class.

Correcting documents

- We're not interested in interactive spelling correction of documents (e.g., MS Word) in this class.
- In IR, we use document correction primarily for OCR'ed documents.

Correcting documents

- We're not interested in interactive spelling correction of documents (e.g., MS Word) in this class.
- In IR, we use document correction primarily for OCR'ed documents.
- The general philosophy in IR is: don't change the documents.

Correcting queries

- First: isolated word spelling correction

Correcting queries

- First: isolated word spelling correction
- Fundamental premise 1: There is a list of “correct words” from which the correct spellings come.

Correcting queries

- First: isolated word spelling correction
- Fundamental premise 1: There is a list of “correct words” from which the correct spellings come.
- Fundamental premise 2: We have a way of computing the **distance** between a misspelled word and a correct word.

Correcting queries

- First: isolated word spelling correction
- Fundamental premise 1: There is a list of “correct words” from which the correct spellings come.
- Fundamental premise 2: We have a way of computing the **distance** between a misspelled word and a correct word.
- Simple spelling correction algorithm: return the “correct” word that has the smallest distance to the misspelled word.

Correcting queries

- First: isolated word spelling correction
- Fundamental premise 1: There is a list of “correct words” from which the correct spellings come.
- Fundamental premise 2: We have a way of computing the **distance** between a misspelled word and a correct word.
- Simple spelling correction algorithm: return the “correct” word that has the smallest distance to the misspelled word.
- Example: *informaton* → *information*

Correcting queries

- First: isolated word spelling correction
- Fundamental premise 1: There is a list of “correct words” from which the correct spellings come.
- Fundamental premise 2: We have a way of computing the **distance** between a misspelled word and a correct word.
- Simple spelling correction algorithm: return the “correct” word that has the smallest distance to the misspelled word.
- Example: *informaton* \rightarrow *information*
- We can use the term vocabulary of the inverted index as the list of correct words.

Correcting queries

- First: isolated word spelling correction
- Fundamental premise 1: There is a list of “correct words” from which the correct spellings come.
- Fundamental premise 2: We have a way of computing the **distance** between a misspelled word and a correct word.
- Simple spelling correction algorithm: return the “correct” word that has the smallest distance to the misspelled word.
- Example: *informaton* \rightarrow *information*
- We can use the term vocabulary of the inverted index as the list of correct words.
- **Why is this problematic?**

Alternatives to using the term vocabulary

- A standard dictionary (Webster's, OED etc.)

Alternatives to using the term vocabulary

- A standard dictionary (Webster's, OED etc.)
- An industry-specific dictionary (for specialized IR systems)

Alternatives to using the term vocabulary

- A standard dictionary (Webster's, OED etc.)
- An industry-specific dictionary (for specialized IR systems)
- The term vocabulary of the collection, appropriately weighted

Distance between misspelled word and “correct” word

- We will study several alternatives.

Distance between misspelled word and “correct” word

- We will study several alternatives.
- Edit distance

Distance between misspelled word and “correct” word

- We will study several alternatives.
- Edit distance
- Levenshtein distance

Distance between misspelled word and “correct” word

- We will study several alternatives.
- Edit distance
- Levenshtein distance
- Weighted edit distance

Distance between misspelled word and “correct” word

- We will study several alternatives.
- Edit distance
- Levenshtein distance
- Weighted edit distance
- k -gram overlap

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- Levenshtein distance: The admissible basic operations are insert, delete, and replace

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- Levenshtein distance: The admissible basic operations are insert, delete, and replace
- Levenshtein distance *dog-do*: 1

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- Levenshtein distance: The admissible basic operations are insert, delete, and replace
- Levenshtein distance *dog-do*: 1
- Levenshtein distance *cat-cart*: 1

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- Levenshtein distance: The admissible basic operations are insert, delete, and replace
- Levenshtein distance *dog-do*: 1
- Levenshtein distance *cat-cart*: 1
- Levenshtein distance *cat-cut*: 1

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- Levenshtein distance: The admissible basic operations are insert, delete, and replace
- Levenshtein distance *dog-do*: 1
- Levenshtein distance *cat-cart*: 1
- Levenshtein distance *cat-cut*: 1
- Levenshtein distance *cat-act*: 2

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- Levenshtein distance: The admissible basic operations are insert, delete, and replace
- Levenshtein distance *dog-do*: 1
- Levenshtein distance *cat-cart*: 1
- Levenshtein distance *cat-cut*: 1
- Levenshtein distance *cat-act*: 2
- Damerau-Levenshtein distance *cat-act*: 1

Edit distance

- The edit distance between string s_1 and string s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- Levenshtein distance: The admissible basic operations are insert, delete, and replace
- Levenshtein distance *dog-do*: 1
- Levenshtein distance *cat-cart*: 1
- Levenshtein distance *cat-cut*: 1
- Levenshtein distance *cat-act*: 2
- Damerau-Levenshtein distance *cat-act*: 1
- Damerau-Levenshtein includes transposition as a fourth possible operation.

Levenshtein distance: Computation

		f	a	s	t
	0	1	2	3	4
c	1	1	2	3	4
a	2	2	1	2	3
t	3	3	2	2	2
s	4	4	3	2	3

Levenshtein distance: algorithm

LEVENSHTEINDISTANCE(s_1, s_2)

```

1  for  $i \leftarrow 0$  to  $|s_1|$ 
2  do  $m[i, 0] = i$ 
3  for  $j \leftarrow 0$  to  $|s_2|$ 
4  do  $m[0, j] = j$ 
5  for  $i \leftarrow 1$  to  $|s_1|$ 
6  do for  $j \leftarrow 1$  to  $|s_2|$ 
7      do if  $s_1[i] = s_2[j]$ 
8          then  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1]\}$ 
9          else  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1] + 1\}$ 
10 return  $m[|s_1|, |s_2|]$ 

```

Operations: insert, delete, replace, copy

Levenshtein distance: algorithm

LEVENSHTEINDISTANCE(s_1, s_2)

```

1  for  $i \leftarrow 0$  to  $|s_1|$ 
2  do  $m[i, 0] = i$ 
3  for  $j \leftarrow 0$  to  $|s_2|$ 
4  do  $m[0, j] = j$ 
5  for  $i \leftarrow 1$  to  $|s_1|$ 
6  do for  $j \leftarrow 1$  to  $|s_2|$ 
7      do if  $s_1[i] = s_2[j]$ 
8          then  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1]\}$ 
9          else  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1] + 1\}$ 
10 return  $m[|s_1|, |s_2|]$ 

```

Operations: insert, delete, replace, copy

Levenshtein distance: algorithm

LEVENSHTEINDISTANCE(s_1, s_2)

```

1  for  $i \leftarrow 0$  to  $|s_1|$ 
2  do  $m[i, 0] = i$ 
3  for  $j \leftarrow 0$  to  $|s_2|$ 
4  do  $m[0, j] = j$ 
5  for  $i \leftarrow 1$  to  $|s_1|$ 
6  do for  $j \leftarrow 1$  to  $|s_2|$ 
7      do if  $s_1[i] = s_2[j]$ 
8          then  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1]\}$ 
9          else  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1] + 1\}$ 
10 return  $m[|s_1|, |s_2|]$ 

```

Operations: insert, delete, replace, copy

Levenshtein distance: algorithm

LEVENSHTEINDISTANCE(s_1, s_2)

```

1  for  $i \leftarrow 0$  to  $|s_1|$ 
2  do  $m[i, 0] = i$ 
3  for  $j \leftarrow 0$  to  $|s_2|$ 
4  do  $m[0, j] = j$ 
5  for  $i \leftarrow 1$  to  $|s_1|$ 
6  do for  $j \leftarrow 1$  to  $|s_2|$ 
7      do if  $s_1[i] = s_2[j]$ 
8          then  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1]\}$ 
9          else  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1] + 1\}$ 
10 return  $m[|s_1|, |s_2|]$ 

```

Operations: insert, delete, **replace**, copy

Levenshtein distance: algorithm

LEVENSHTEINDISTANCE(s_1, s_2)

```

1  for  $i \leftarrow 0$  to  $|s_1|$ 
2  do  $m[i, 0] = i$ 
3  for  $j \leftarrow 0$  to  $|s_2|$ 
4  do  $m[0, j] = j$ 
5  for  $i \leftarrow 1$  to  $|s_1|$ 
6  do for  $j \leftarrow 1$  to  $|s_2|$ 
7      do if  $s_1[i] = s_2[j]$ 
8          then  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1]\}$ 
9          else  $m[i, j] = \min\{m[i-1, j] + 1, m[i, j-1] + 1, m[i-1, j-1] + 1\}$ 
10 return  $m[|s_1|, |s_2|]$ 

```

Operations: insert, delete, replace, copy

Levenshtein distance: Example

			f		a		s		t	
		<u>0</u>	<u>1</u>	<u>1</u>	<u>2</u>	<u>2</u>	<u>3</u>	<u>3</u>	<u>4</u>	<u>4</u>
c		<u>1</u>	<u>1</u>	<u>2</u>	<u>2</u>	<u>3</u>	<u>3</u>	<u>4</u>	<u>4</u>	<u>5</u>
		<u>1</u>	<u>2</u>	<u>1</u>	<u>2</u>	<u>2</u>	<u>3</u>	<u>3</u>	<u>4</u>	<u>4</u>
a		<u>2</u>	<u>2</u>	<u>2</u>	<u>1</u>	<u>3</u>	<u>3</u>	<u>4</u>	<u>4</u>	<u>5</u>
		<u>2</u>	<u>3</u>	<u>2</u>	<u>3</u>	<u>1</u>	<u>2</u>	<u>2</u>	<u>3</u>	<u>3</u>
t		<u>3</u>	<u>3</u>	<u>3</u>	<u>3</u>	<u>2</u>	<u>2</u>	<u>3</u>	<u>2</u>	<u>4</u>
		<u>3</u>	<u>4</u>	<u>3</u>	<u>4</u>	<u>2</u>	<u>3</u>	<u>2</u>	<u>3</u>	<u>2</u>
s		<u>4</u>	<u>4</u>	<u>4</u>	<u>4</u>	<u>3</u>	<u>2</u>	<u>3</u>	<u>3</u>	<u>3</u>
		<u>4</u>	<u>5</u>	<u>4</u>	<u>5</u>	<u>3</u>	<u>4</u>	<u>2</u>	<u>3</u>	<u>3</u>

Each cell of Levenshtein matrix

cost of getting here from my upper left neighbor (copy or replace)	cost of getting here from my upper neighbor (delete)
cost of getting here from my left neighbor (insert)	the minimum of the three possible “move- ments”; the cheapest way of getting here

Dynamic programming (Cormen et al.)

- Optimal substructure: The optimal solution to the problem contains within it optimal solutions to subproblems.

Dynamic programming (Cormen et al.)

- Optimal substructure: The optimal solution to the problem contains within it optimal solutions to subproblems.
- Overlapping subproblems: The optimal solutions to subproblems (“subsolutions”) overlap. These subsolutions are computed over and over again when computing the global optimal solution.

Dynamic programming (Cormen et al.)

- Optimal substructure: The optimal solution to the problem contains within it optimal solutions to subproblems.
- Overlapping subproblems: The optimal solutions to subproblems (“subsolutions”) overlap. These subsolutions are computed over and over again when computing the global optimal solution.
- Optimal substructure: We compute minimum distance of substrings in order to compute the minimum distance of the entire string.

Dynamic programming (Cormen et al.)

- Optimal substructure: The optimal solution to the problem contains within it optimal solutions to subproblems.
- Overlapping subproblems: The optimal solutions to subproblems (“subsolutions”) overlap. These subsolutions are computed over and over again when computing the global optimal solution.
- Optimal substructure: We compute minimum distance of substrings in order to compute the minimum distance of the entire string.
- Overlapping subproblems: Need most distances of substrings 3 times (moving right, diagonally, down)

<http://ifnlp.org/lehre/teaching/2008-SS/ir/editdist2.pdf>

Exercise

- Given: *cat* and *catcat*

Exercise

- Given: *cat* and *catcat*
- Compute the matrix of Levenshtein distances

Exercise

- Given: *cat* and *catcat*
- Compute the matrix of Levenshtein distances
- Read out the editing operations that transform *cat* into *catcat*

Weighted edit distance

- As above, but weight of an operation depends on the characters involved.

Weighted edit distance

- As above, but weight of an operation depends on the characters involved.
- Meant to capture keyboard errors, e.g., m more likely to be mistyped as n than as q .

Weighted edit distance

- As above, but weight of an operation depends on the characters involved.
- Meant to capture keyboard errors, e.g., m more likely to be mistyped as n than as q .
- Therefore, replacing m by n is a smaller edit distance than by q .

Weighted edit distance

- As above, but weight of an operation depends on the characters involved.
- Meant to capture keyboard errors, e.g., m more likely to be mistyped as n than as q .
- Therefore, replacing m by n is a smaller edit distance than by q .
- We now require a weight matrix as input.

Weighted edit distance

- As above, but weight of an operation depends on the characters involved.
- Meant to capture keyboard errors, e.g., m more likely to be mistyped as n than as q .
- Therefore, replacing m by n is a smaller edit distance than by q .
- We now require a weight matrix as input.
- Modify dynamic programming to handle weights.

Using edit distance

- Given query, first enumerate all character sequences within a preset (possibly weighted) edit distance

Using edit distance

- Given query, first enumerate all character sequences within a preset (possibly weighted) edit distance
- Intersect this set with list of “correct” words

Using edit distance

- Given query, first enumerate all character sequences within a preset (possibly weighted) edit distance
- Intersect this set with list of “correct” words
- Then suggest terms you found to the user.

Using edit distance

- Given query, first enumerate all character sequences within a preset (possibly weighted) edit distance
- Intersect this set with list of “correct” words
- Then suggest terms you found to the user.
- Or do automatic correction – but this is potentially expensive and disempowers the user.

k -gram indexes for spelling correction

- Enumerate all k -grams in the query term

k -gram indexes for spelling correction

- Enumerate all k -grams in the query term
- Use the k -gram index to retrieve “correct” words that match query term k -grams

k -gram indexes for spelling correction

- Enumerate all k -grams in the query term
- Use the k -gram index to retrieve “correct” words that match query term k -grams
- Threshold by number of matching k -grams

k -gram indexes for spelling correction

- Enumerate all k -grams in the query term
- Use the k -gram index to retrieve “correct” words that match query term k -grams
- Threshold by number of matching k -grams
- E.g., only vocabulary terms that differ by at most 3 k -grams


k -gram indexes for spelling correction

- Enumerate all k -grams in the query term
- Use the k -gram index to retrieve “correct” words that match query term k -grams
- Threshold by number of matching k -grams
- E.g., only vocabulary terms that differ by at most 3 k -grams
- Example: bigram index, misspelled word *bordroom*

k -gram indexes for spelling correction

- Enumerate all k -grams in the query term
- Use the k -gram index to retrieve “correct” words that match query term k -grams
- Threshold by number of matching k -grams
- E.g., only vocabulary terms that differ by at most 3 k -grams
- Example: bigram index, misspelled word *bordroom*
- Bigrams: *bo, or, rd, dr, ro, oo, om*

k-gram indexes for spelling correction: *bordroom*



Example with trigrams

- Issue: Fixed number of k -grams that differ does not work for words of differing length.

Example with trigrams

- Issue: Fixed number of k -grams that differ does not work for words of differing length.
- Suppose the correct word is NOVEMBER

Example with trigrams

- Issue: Fixed number of k -grams that differ does not work for words of differing length.
- Suppose the correct word is NOVEMBER
- Trigrams: *nov*, *ove*, *vem*, *emb*, *mbe*, *ber*

Example with trigrams

- Issue: Fixed number of k -grams that differ does not work for words of differing length.
- Suppose the correct word is NOVEMBER
- Trigrams: *nov*, *ove*, *vem*, *emb*, *mbe*, *ber*
- And the query term is DECEMBER

Example with trigrams

- Issue: Fixed number of k -grams that differ does not work for words of differing length.
- Suppose the correct word is NOVEMBER
- Trigrams: *nov, ove, vem, emb, mbe, ber*
- And the query term is DECEMBER
- Trigrams: *dec, ece, cem, emb, mbe, ber*

Example with trigrams

- Issue: Fixed number of k -grams that differ does not work for words of differing length.
- Suppose the correct word is NOVEMBER
- Trigrams: *nov, ove, vem, emb, mbe, ber*
- And the query term is DECEMBER
- Trigrams: *dec, ece, cem, emb, mbe, ber*
- So 3 trigrams overlap (out of 6 in each term)

Example with trigrams

- Issue: Fixed number of k -grams that differ does not work for words of differing length.
- Suppose the correct word is NOVEMBER
- Trigrams: *nov, ove, vem, emb, mbe, ber*
- And the query term is DECEMBER
- Trigrams: *dec, ece, cem, emb, mbe, ber*
- So 3 trigrams overlap (out of 6 in each term)
- How can we turn this into a normalized measure of overlap?

Jaccard coefficient

- A commonly used measure of overlap of two sets

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets
- Jaccard coefficient:

$$\frac{|A \cap B|}{|A \cup B|}$$

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets
- Jaccard coefficient:

$$\frac{|A \cap B|}{|A \cup B|}$$

- Values if A and B have the same elements? If they are disjoint?

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets
- Jaccard coefficient:

$$\frac{|A \cap B|}{|A \cup B|}$$

- Values if A and B have the same elements? If they are disjoint?
- A and B don't have to be the same size.

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets
- Jaccard coefficient:

$$\frac{|A \cap B|}{|A \cup B|}$$

- Values if A and B have the same elements? If they are disjoint?
- A and B don't have to be the same size.
- Always assigns a number between 0 and 1.

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets
- Jaccard coefficient:

$$\frac{|A \cap B|}{|A \cup B|}$$

- Values if A and B have the same elements? If they are disjoint?
- A and B don't have to be the same size.
- Always assigns a number between 0 and 1.
- december/november example: Jaccard coefficient?

Jaccard coefficient

- A commonly used measure of overlap of two sets
- Let A and B be two sets
- Jaccard coefficient:

$$\frac{|A \cap B|}{|A \cup B|}$$

- Values if A and B have the same elements? If they are disjoint?
- A and B don't have to be the same size.
- Always assigns a number between 0 and 1.
- december/november example: Jaccard coefficient?
- Application to spelling correction: declare a match if the coefficient is, say, > 0.8 .

Context-sensitive spelling correction

- Our example was: *an asteroid that fell **form** the sky*

Context-sensitive spelling correction

- Our example was: *an asteroid that fell form the sky*
- How can we correct *form* here?

Context-sensitive spelling correction

- Our example was: *an asteroid that fell form the sky*
- How can we correct *form* here?
- Ideas?

Context-sensitive spelling correction

- Our example was: *an asteroid that fell form the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction

Context-sensitive spelling correction

- Our example was: *an asteroid that fell **form** the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term

Context-sensitive spelling correction

- Our example was: *an asteroid that fell form the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term
 - for *flew form munich*: *flea* for *flew*, *from* for *form*, *munch* for *munich*

Context-sensitive spelling correction

- Our example was: *an asteroid that fell form the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term
 - for *flew form munich*: *flea* for *flew*, *from* for *form*, *munch* for *munich*
 - Now try all possible resulting phrases as queries with one word “fixed” at a time

Context-sensitive spelling correction

- Our example was: *an asteroid that fell form the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term
 - for *flew form munich*: *flea* for *flew*, *from* for *form*, *munch* for *munich*
 - Now try all possible resulting phrases as queries with one word “fixed” at a time
 - Try query “*flea form munich*”

Context-sensitive spelling correction

- Our example was: *an asteroid that fell form the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term
 - for *flew form munich*: *flea* for *flew*, *from* for *form*, *munch* for *munich*
 - Now try all possible resulting phrases as queries with one word “fixed” at a time
 - Try query “*flea form munich*”
 - Try query “*flew from munich*”

Context-sensitive spelling correction

- Our example was: *an asteroid that fell **form** the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term
 - for *flew form munch*: *flea* for *flew*, *from* for *form*, *munch* for *munch*
 - Now try all possible resulting phrases as queries with one word “fixed” at a time
 - Try query “*flea form munch*”
 - Try query “*flew from munch*”
 - Try query “*flew form munch*”

Context-sensitive spelling correction

- Our example was: *an asteroid that fell **form** the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term
 - for *flew form munch*: *flea* for *flew*, *from* for *form*, *munch* for *munch*
 - Now try all possible resulting phrases as queries with one word “fixed” at a time
 - Try query “*flea form munch*”
 - Try query “*flew from munch*”
 - Try query “*flew form munch*”
 - The correct query “*flew from munch*” has the most hits.

Context-sensitive spelling correction

- Our example was: *an asteroid that fell **form** the sky*
- How can we correct *form* here?
- Ideas?
- One idea: **hit-based** spelling correction
 - Retrieve “correct” terms close to each query term
 - for *flew form munich*: *flea* for *flew*, *from* for *form*, *munch* for *munich*
 - Now try all possible resulting phrases as queries with one word “fixed” at a time
 - Try query “*flea form munich*”
 - Try query “*flew from munich*”
 - Try query “*flew form munch*”
 - The correct query “*flew from munich*” has the most hits.
- Suppose we have 7 alternatives for *flew*, 19 for *form* and 3 for *munich*, how many “corrected” phrases will we enumerate?

Context-sensitive spelling correction

- The “hit-based” algorithm we just outlined is not very efficient.

Context-sensitive spelling correction

- The “hit-based” algorithm we just outlined is not very efficient.
- More efficient alternative: look at “collection” of queries, not documents

General issues in spelling correction

- User interface

General issues in spelling correction

- User interface
 - automatic vs. suggested correction

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.
 - What about multiple possible corrections?

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.
 - What about multiple possible corrections?
 - Tradeoff: simple vs. powerful UI

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.
 - What about multiple possible corrections?
 - Tradeoff: simple vs. powerful UI
- Cost

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.
 - What about multiple possible corrections?
 - Tradeoff: simple vs. powerful UI
- Cost
 - Spelling correction is potentially expensive.

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.
 - What about multiple possible corrections?
 - Tradeoff: simple vs. powerful UI
- Cost
 - Spelling correction is potentially expensive.
 - Avoid running on every query?

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.
 - What about multiple possible corrections?
 - Tradeoff: simple vs. powerful UI
- Cost
 - Spelling correction is potentially expensive.
 - Avoid running on every query?
 - Maybe just on queries that match few documents.

General issues in spelling correction

- User interface
 - automatic vs. suggested correction
 - *Did you mean* only works for one suggestion.
 - What about multiple possible corrections?
 - Tradeoff: simple vs. powerful UI
- Cost
 - Spelling correction is potentially expensive.
 - Avoid running on every query?
 - Maybe just on queries that match few documents.
 - Guess: Spelling correction of major search engines is efficient enough to be run on every query.

Peter Norvig's complete spelling corrector in only 21 lines of code!

Outline

- 1 Recap
- 2 Dictionaries
- 3 Wildcard queries
- 4 Spelling correction
- 5 Soundex**

Soundex

- Soundex is the basis for finding **phonetic** (as opposed to orthographic) alternatives.

Soundex

- Soundex is the basis for finding **phonetic** (as opposed to orthographic) alternatives.
- Example: *chebyshev* / *tchebyscheff*

Soundex

- Soundex is the basis for finding **phonetic** (as opposed to orthographic) alternatives.
- Example: *chebyshev* / *tchebyscheff*
- Algorithm:

Soundex

- Soundex is the basis for finding **phonetic** (as opposed to orthographic) alternatives.
- Example: *chebyshev* / *tchebyscheff*
- Algorithm:
 - Turn every token to be indexed into a 4-character reduced form

Soundex

- Soundex is the basis for finding **phonetic** (as opposed to orthographic) alternatives.
- Example: *chebyshev* / *tchebyscheff*
- Algorithm:
 - Turn every token to be indexed into a 4-character reduced form
 - Do the same with query terms

Soundex

- Soundex is the basis for finding **phonetic** (as opposed to orthographic) alternatives.
- Example: *chebyshev* / *tchebyscheff*
- Algorithm:
 - Turn every token to be indexed into a 4-character reduced form
 - Do the same with query terms
 - Build and search an index on the reduced forms

Soundex algorithm

- 1 Retain the first letter of the term.

Soundex algorithm

- 1 Retain the first letter of the term.
- 2 Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'

Soundex algorithm

- 1 Retain the first letter of the term.
- 2 Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- 3 Change letters to digits as follows:

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1
 - C, G, J, K, Q, S, X, Z to 2

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1
 - C, G, J, K, Q, S, X, Z to 2
 - D, T to 3

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1
 - C, G, J, K, Q, S, X, Z to 2
 - D, T to 3
 - L to 4

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1
 - C, G, J, K, Q, S, X, Z to 2
 - D, T to 3
 - L to 4
 - M, N to 5

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1
 - C, G, J, K, Q, S, X, Z to 2
 - D, T to 3
 - L to 4
 - M, N to 5
 - R to 6

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1
 - C, G, J, K, Q, S, X, Z to 2
 - D, T to 3
 - L to 4
 - M, N to 5
 - R to 6
- ➍ Repeatedly remove one out of each pair of consecutive identical digits

Soundex algorithm

- ➊ Retain the first letter of the term.
- ➋ Change all occurrences of the following letters to '0' (zero): 'A', 'E', 'I', 'O', 'U', 'H', 'W', 'Y'
- ➌ Change letters to digits as follows:
 - B, F, P, V to 1
 - C, G, J, K, Q, S, X, Z to 2
 - D, T to 3
 - L to 4
 - M, N to 5
 - R to 6
- ➍ Repeatedly remove one out of each pair of consecutive identical digits
- ➎ Remove all zeros from the resulting string; pad the resulting string with trailing zeros and return the first four positions, which will consist of a letter followed by three digits

Example: Soundex of *HERMAN*

- Retain H

Example: Soundex of *HERMAN*

- Retain H
- *ERMAN* \rightarrow *ORMON*

Example: Soundex of *HERMAN*

- Retain H
- *ERMAN* \rightarrow *ORM0N*
- *ORM0N* \rightarrow *06505*

Example: Soundex of *HERMAN*

- Retain H
- *ERMAN* \rightarrow *ORM0N*
- *ORM0N* \rightarrow *06505*
- *06505* \rightarrow *06505*

Example: Soundex of *HERMAN*

- Retain H
- *ERMAN* \rightarrow *ORM0N*
- *ORM0N* \rightarrow *06505*
- *06505* \rightarrow *06505*
- *06505* \rightarrow *655*

Example: Soundex of *HERMAN*

- Retain H
- *ERMAN* → *ORM0N*
- *ORM0N* → *06505*
- *06505* → *06505*
- *06505* → *655*
- Return *H655*

Example: Soundex of *HERMAN*

- Retain H
- *ERMAN* \rightarrow *ORM0N*
- *ORM0N* \rightarrow *06505*
- *06505* \rightarrow *06505*
- *06505* \rightarrow *655*
- Return *H655*
- Will *HERMANN* generate the same code?

Compute soundex code of your last name.

How useful is Soundex?

- Not very – for information retrieval

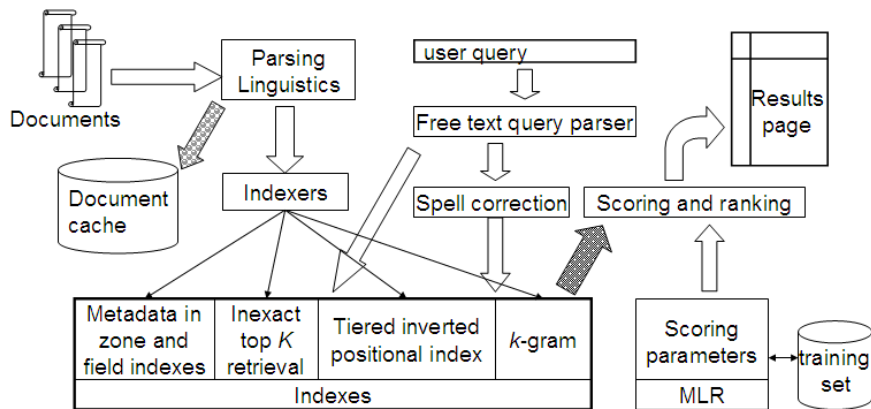
How useful is Soundex?

- Not very – for information retrieval
- Ok for “high recall” tasks in other applications (e.g., Interpol)

How useful is Soundex?

- Not very – for information retrieval
- Ok for “high recall” tasks in other applications (e.g., Interpol)
- Zobel and Dart (1996) suggest better alternatives for phonetic matching in IR.

The complete search system



Resources

- Chapter 3 of IIR

Resources

- Chapter 3 of IIR
- Resources at <http://ifnlp.org/ir>

Resources

- Chapter 3 of IIR
- Resources at <http://ifnlp.org/ir>
- Soundex demo

Resources

- Chapter 3 of IIR
- Resources at <http://ifnlp.org/ir>
- Soundex demo
- Levenshtein distance demo

Resources

- Chapter 3 of IIR
- Resources at <http://ifnlp.org/ir>
- Soundex demo
- Levenshtein distance demo
- Levenshtein distance slides

Resources

- Chapter 3 of IIR
- Resources at <http://ifnlp.org/ir>
- Soundex demo
- Levenshtein distance demo
- Levenshtein distance slides
- Peter Norvig's spelling corrector