

Dynamic Programming Problem Set Solution

Amir Ebrahimifard - Pourya Jahandideh

1. (a) The algorithm written in C:

```
int T(int n) {
    int sum = 0;
    if (n==0 || n==1)
        return 2;
    for(int j=1; j<=n-1; j++)
        sum+=T(j)*T(j-1);
    return sum;
}
```

Let $N(n)$ denote the number of operations required to calculate $T(n)$. Note that $N(2) = 2$. So if n is even, from the algorithm above we have:

$$\begin{aligned} N(n) &= \sum_{i=1}^{n-1} 2 * N(i) - N(n-1) + N(0) + 2 * (n-1) \\ &\geq N(n-1) + N(n-2) \geq 2 * N(n-2) \\ &\geq 2 * 2 * N(n-4) \geq \dots \geq 2^{\frac{n}{2}}. \end{aligned}$$

When n is odd, it's the same idea. Hence it's exponential.

- (b) An algorithm in C:

```
T[0]=T[1]=2;
for(j=2; j<=n; j++)
{
    T[j] = 0;
    for(k=1; k<j; k++)
        T[j] += T[k]*T[k-1];
}
return T[n];
```

- (c) An algorithm in C:

```
T[0]=T[1]=2;
T[2]=T[0]*T[1];
for(j=3; j<=n; j++)
    T[j]=T[j-1] + T[j-1]*T[j-2];
return T[n];
```

2. The solution to the problem of the shortest common subsequence for three strings is essentially identical to the solution with 2 strings. By treating the $T(i, j, k)$'s as array entries and updating the table in the appropriate way we can get the following $O(n^3)$ time algorithm. Let a be the length of A , b be the length of B , and c be the length of C .

```

for j = 0 to b
  for k = 0 to c
    T(0,j,k)=0

for k = 0 to c
  for i = 0 to a
    T(i,0,k)=0

for i = 0 to a
  for j = 0 to b
    T(i,j,0)=0

for i = 1 to a do
  for j = 1 to b do
    for k = 1 to c do
      if a(i) = b(j) = c(k) then
        T(i,j,k)=T(i-1,j-1,k-1) + 1
      else
        T(i,j,k)= MAX(T(i, j-1, k), T(i-1, j, k), T(i, j, k-1))

```

3. We present an algorithm to compute the shortest common super-sequence of two strings: Let $A = a_1 \dots a_m$ and $B = b_1 \dots b_n$. Let $M[i, j]$ denote the length of the shortest common super-sequence of string $a_1 \dots a_i, b_1 \dots b_j$, for $1 \leq i \leq m$ and $1 \leq j \leq n$. Note that the last letter in the shortest super-sequence of a_1, \dots, a_i and b_1, \dots, b_j is either a_i or b_j . If the last letter is a_i then the previous letters are the solution to the subproblem for a_1, \dots, a_{i-1} and b_1, \dots, b_j . If the last letter is b_j then the previous letters are the solution to the subproblem for a_1, \dots, a_i and b_1, \dots, b_{j-1} . Hence, $M(i, j) =$

$$\min(M(i-1, j-1) + 1 \text{ if } a_i = b_j, M(i-1, j) + 1, M(i, j-1) + 1).$$

This table update can be embedded in two nested loops, the first that goes from $i = 1$ to m and the second with goes from $j = 1$ to n . This gives an $\Theta(n^2)$ time algorithm. Note that the actual sequence can be computed by following the updates back from $M(m, n)$. Code follows:

```

for i = 1 to m do
  for j = 1 to n do
    if  $a_i = b_j$  then
       $M[i, j] = M[i-1, j-1] + 1$ 
    else  $M[i, j] = \min(M[i-1, j], M[i, j-1]) + 1$ 

```

4. [This solution was adapted from Brian Wongchaowart's homework writeup.]

- (a) The path taken by the algorithm through the table to reconstruct the shortest common super-sequence is highlighted in bold.

Let M be the name of our table, let $A = \text{zxyyzz}$ and $B = \text{zzyxzy}$.

M		B						
j=		0	1	2	3	4	5	6
			z	z	y	x	z	y
i=0		0	1	2	3	4	5	6
1	z	1	1	2	3	4	5	6
2	x	2	2	3	4	4	5	6
A 3	y	3	3	4	4	5	6	6
4	y	4	4	5	5	6	7	7
5	z	5	5	5	6	7	7	8
6	z	6	6	6	7	8	8	9

Following this path from $M[6,6]$ to $M[0,0]$, the shortest common super-sequence of A and B is zzyxzzyzz . (For details see part c.)

- (b) If the strings A and B have length m and n , respectively, the table entry $M[m, n]$ gives the length of the shortest common super-sequence of A and B . In this example, the bottom-right entry in the table is 9, which is the length of the shortest common super-sequence of zxyyzz and zzyxzy .
- (c) The letters of string A label the rows of the table, and the letters of string B label the columns. The shortest common super-sequence is constructed in reverse order starting from the bottom-right entry. A letter is added to the beginning of the partial super-sequence constructed so far as one moves past its row or column, so moving left adds the letter at the top of the current column, and moving up adds the letter at the head of the current row. A diagonal move is only allowed when the current row and column are marked with the same letter; that letter is added once to the common super-sequence but is accounted for (moved over) in both words.

The idea is to “reverse” the code that built the table, eventually reaching the upper-left corner (the entry containing the 0) in as few moves as possible, at which point every letter of both words will have been added to the super-sequence. Note that in the first row and column there is obviously only one possible direction in which to move. Some pseudocode for the trace back:

```

i = m, j = n, S = ""
while M[i, j] > 0 //do until we've reached upper left corner of M
  if i == 0 //we're in the first row
    S = b_j + S; //add the letter from B
  
```

```

        j - -; //move left one column
    else if j == 0 //we're in the first column
        S = a_i + S; //add letter from A
        i - -; //move up one row
    else if a_i == b_j then
        S = a_i + S; //add this letter (it's in both A and B)
        i - -, j - -; //move diagonally up and left
    else if M(i - 1, j) ≤ M(i, j - 1) then
        S = a_i + S; //add letter from A
        i - -; //move up one row
    else
        S = b_j + S; //add letter from B
        j - -; //move left one column
    end if
end while
return S

```

5. We present an algorithm to compute the minimum edit distance of two strings. Note that:

- (a) If it were possible to convert a_1, \dots, a_{m-1} into b_1, \dots, b_n , one could complete the transformation of A into B by deleting a_m .
- (b) If it were possible to convert a_1, \dots, a_m into b_1, \dots, b_{n-1} , one could complete the transformation by adding b_n to A .
- (c) If it were possible to convert a_1, \dots, a_{m-1} into b_1, \dots, b_{n-1} , one could complete the transformation by replacing a_m with b_n .
- (d) If string A is empty and B is not empty then the conversion can only be done by inserting all remaining characters of B into A .
- (e) If string B is empty and A is not empty then the conversion can only be done by deleting all remaining characters in A (assuming we can't replace a character with the empty character).

Now let $A[i, j]$ be the minimum cost of transforming a_1, \dots, a_i into b_1, \dots, b_j . The algorithm is:

MinumumEditDistance(A, B)

$A[0, 0] = 0$

for $i = 1$ to m

$A[i, 0] = i * 3$

for $j = 1$ to n

$A[0, j] = j * 4$

```

for  $i = 1$  to  $m$ 
  for  $j = 1$  to  $n$ 
    if  $a_i = b_j$  then
       $A[i, j] = A[i - 1, j - 1]$ 
    else  $A[i, j] = \min(A[i - 1, j] + 3, A[i, j - 1] + 4, A[i - 1, j - 1] + 5)$ 

```

Starting from $A[m, n]$, we can trace backwards through the table to determine which operations were performed at each step.

6. [This solution is courtesy Matthias Grabmair (in collaboration with Ian Wong).]

For K_1, K_2, K_3, K_4, K_5 , our algorithm produces the following table of expected access times

	0	1	2	3	4	5
1	0	.5	.6	.85	1.4	2.15
2		0	.05	.2	.55	1.05
3			0	.1	.4	.9
4				0	.2	.65
5					0	.25

Given this table, we can produce the table of roots that correspond to the above access times:


	0	1	2	3	4	5
1	0	1	1	1	1	1
2		0	2	3	4	4
3			0	3	4	4
4				0	4	5
5					0	5

Hence, the tree looks as follows:

```

1
 \

```



The expected access time of the tree is found in the upper right corner of the table, 2.15. The tree can be constructed from the table above as follows. Start with the upper right corner (entry [2,5]) in the table of roots to see that node 1 is at the root of the tree. Hence the remaining nodes 2 through 4 are all in one subtree that is the right child of the root. The subtree has minimal weight when the root is 4, given in position [2,5] of the table. Since node 4 is the new root node, 5 inevitably becomes 4's right child node. For the subtree that is the left child of node 4, we need to figure out which of 2 or 3 is the root. According to table cell [2,3] it is node 3.

7. We use dynamic programming to formulate an algorithm to compute the minimum polygon triangulation. Note that after the first "cut" is made, the original polygon P will be divided into two new polygons P' and P'' . The triangulations P' and P'' are completely independent of each other. The cheapest way to triangulate P is the minimum of:
 - (a) For $3 \leq k \leq n-2$, the length of (v_1, v_k) plus the length of (v_k, v_n) plus the cheapest way to triangulate $P' = \{v_1, \dots, v_k\}$ plus the cheapest way to triangulate $P'' = \{v_k, \dots, v_n\}$.
 - (b) The length of (v_2, v_n) plus the cheapest way to triangulate $P' = \{v_2, \dots, v_n\}$.
 - (c) The length of (v_1, v_{n-1}) plus the cheapest way to triangulate $P' = \{v_1, \dots, v_{n-1}\}$.

Create a two-dimensional array A . $A[i, j]$ = minimum weight of triangulation of the polygon defined by v_i, \dots, v_j . Note we can initialize the diagonal (where $i = j$) to zeroes and will not use any entries below the diagonal (where $i > j$).

for $i = n$ **downto** 1 **do**
 for $j = i + 1$ **to** n **do**

$$T[i, j] = \min(\min_{i+2 \leq k \leq j-2} (d(i, k) + d(k, j) + A[i, k] + A[k, j]), \\ d(j, i+1) + A[i+1, j], \\ d(i, j-1) + A[i, j-1])$$

8. We present an algorithm to compute the maximum consecutive sum of n integers x_1, \dots, x_n . We define two functions:

- $\text{MCS}(i)$ = Maximum Consecutive Sum of the first i integers.
- $\text{MSS}(i)$ = Maximum Consecutive Sum of the first i integers that uses x_i .

$\text{MSS}(i)$ is computed by:

$$\text{MSS}(i) = \text{MAX}(0, \text{MSS}(i-1) + x_i)$$

And $\text{MCS}(n)$ is computed by:

$$\text{MCS}(i) = \text{MAX}(\text{MCS}(i-1), \text{MSS}(i-1) + x_i)$$

By placing the two assignment statements inside a loop where i runs from 1 to n , we get a linear time algorithm.

9. The input to this problem is a tree T with weights on the edges. The goal is to find the path in T with minimum aggregate weight. A path is a collection of adjacent vertices, where no vertex is used more than once.

Root the tree at an arbitrary node r , and process the tree in postorder. We generalize the induction hypothesis to compute not only the shortest path in each subtree, but also the shortest path with one endpoint being the root of the subtree. Consider an arbitrary node v with branches to k descendants w_1, w_2, \dots, w_k . For each such node v the algorithm computes the following information:

- $\text{best}(v)$ = the minimum weight of a path for the subtree rooted at v .
- $\text{root-best}(v)$ = the minimum weight of a path ending at v in the subtree rooted v .

At node v , the algorithm first recursively computes $\text{best}(w_i)$ and $\text{root-best}(w_i)$ for each descendant subtree. It then computes $\text{best}(v)$ and $\text{root-best}(v)$ using the following recurrence relations that correspond to the two cases identified above:

$$\text{rootbest}(v) = \min(0, \min_{i=1}^k (d(v, w_i) + \text{rootbest}(w_i)))$$

$$\text{best}(v) = \min(\text{rootbest}(v), \min_i (\text{best}(w_i)), \min_{i,j} (\text{rootbest}(w_i) + \text{rootbest}(w_j) + d(v, w_i) + d(v, w_j)))$$

Its not too hard (but not completely trivial) to see that this can be implemented in linear time.

10. We present the following dynamic programming algorithm to compute the AVL tree with minimum expected access time. The main idea is to strengthen the inductive hypothesis to compute the AVL tree of each height with the minimum expected access time. Define $A[i, j, h]$ as the minimized expected access time for a tree with height h on the keys K_i, K_{i+1}, \dots, K_j . We then get the following recurrence:

$$A[i, j, h] = \min_{i \leq r \leq j} \left(\begin{aligned} &A[i, r-1, h-1] + A[r+1, j, h-2] + \sum_{a=i}^j p_a, \\ &A[i, r-1, h-1] + A[r+1, j, h-1] + \sum_{a=i}^j p_a, \\ &A[i, r-1, h-2] + A[r+1, j, h-1] + \sum_{a=i}^j p_a \end{aligned} \right)$$

Here r is the root of the new tree, and p_a is the probability of accessing key K_a . Note that if the tree is going to be of height h then one of its two subtrees must be of height $h-1$ and the other can be of height at most $h-1$. Since the tree is an AVL tree the heights of the two subtrees can differ by at most 1. One can then get code by wrapping this assignment statements in a loop for i from n to 1, a loop for j from i to n , and a loop for h from 0 to n (actually you can replace n here by something like $\sqrt{2} \log n$ if you know that AVL trees are balanced). The final minimum expected depth can be found by taking the minimum over all h of $A[1, n, h]$, and the tree can be recovered the same way that it was for the problem on computing the binary search tree with minimum expected access time.

11. We give dynamic programming algorithm to compute the non-overlapping collection of intervals I_1, \dots, I_n with maximum measure. We consider the intervals by increasing order of their left endpoints. The main idea is to strengthen the inductive hypothesis to compute the maximum measure non-overlapping collection of intervals with a particular rightmost interval, for all possible choices of rightmost interval. Define $S[i, j]$ as the total length of the longest non-overlapping interval set among the first i intervals such that the right-most interval in this set is the j th one. Then $S[i, j]$ is the maximum over all intervals I_k , such that the right endpoint of I_k is to the left of the left endpoint of I_j , of $S[k, k]$ plus the length of I_j . By wrapping this in a loop for i from 1 to n , and for j from 1 to i , you get the algorithm. The maximum measure of a non-overlapping collection of intervals can then be found by taking the maximum $S[n, j]$ for all j .
12. No solution given.
13. No solution given.

14. We present a solution for the problem of determining if the values of n items can be added and subtracted in such a way that $\sum_{k=1}^n (-1)^{x_k} v_k = 0$. We use dynamic programming. As pruning rule we remark that if we have two solutions with the same value we only need to keep one. Here we cannot discard solutions with negative values or with values larger than L since subtraction is allowed. However, if we define $L = \sum_{i=1}^n v_i$, we know that any solution will have its value in $\{-L \dots L\}$, and thus we will have at most $2L$ solutions to keep.

For $i = 1 \dots n$ and $j = -L \dots L$ let $A(i, j)$ be a boolean variable that indicates whether or not there is a solution with the first i objects that has a value equal to j . That is, $A(i, j)$ is true if there is a solution to

$$\sum_{k=1}^i (-1)^{x_k} v_k = j$$

where each x_k is either 1 or 0.

With the initialization $A(0, 0) = T$ and $A(0, j) = F$ for $j \neq 0$, we can fill the table row by row using the following rule:

$$\text{if } A(i, j) \text{ then } A(i+1, j+v_{i+1}) = T \text{ and } A(i+1, j-v_{i+1}) = T.$$

That is, if there is a solution for j at level i then there is a solution for $j+v_{i+1}$ and $j-v_{i+1}$ at level $i+1$.

At the end, the entry $A(n, 0)$ will indicate whether there is a solution for the problem or not. If you want you can then construct a solution by tracing a “path” back in the table starting at $A(n, 0)$.

The size of the table is $n \times 2L$. The time to fill an entry is constant. Therefore the total time is $O(nL)$ which is polynomial in $n + L$.

15. We present an algorithm for determining if there is a subset of n items whose total value is $L \bmod n$. We use dynamic programming. As pruning rule we remark that if we have two solutions with the same value modulo n we only need to keep one. Therefore there are at most n solutions to keep at each level.

For $i = 1 \dots n$ and $j = 0 \dots n-1$ let $A(i, j)$ be a boolean variable that indicate whether or not there is a solution with the first i objects that has a value equal to $j \bmod n$. That is, $A(i, j)$ is true if there is a solution to

$$A(i, j) = \left(\sum_{k=1}^i x_k v_k \right) \bmod n = j$$

where each x_k is either 1 or 0.

With the initialization $A(0,0) = T$ and $A(0,j) = F$ for $j = 1 \dots n-1$, we can fill the table row by row using the following recurrence relation:

$$A(i,j) = A(i-1,j) \text{ or } A(i-1, (j - v_i) \bmod n).$$

That is, there is a solution modulo j at level i if there is a solution for j at level $i-1$ or if you can add v_i to a solution at level $i-1$ and get j modulo n .

At the end, the entry $A(n, L \bmod n)$ will indicate whether there is a solution for the problem or not. If you want you can then construct a solution by tracing a “path” back in the table starting at $A(n, L \bmod n)$.

The size of the table is n^2 . The time to fill an entry is constant. Therefore the total time is $O(n^2)$ which is polynomial in $n + L$.

16. We use dynamic programming for the problem of obtaining the maximum value from a subcollection (allowing repetition) of n items, subject to the restriction that the total weight of the set cannot exceed W . As pruning rule we remark that if we have two solutions with the same weight we only need to keep the one with the highest value and that we only need to keep solutions with weight no greater than W .

For $i = 1 \dots n$ and $j = 0 \dots W$ we compute $A(i,j)$ = the maximum value of an assignment of the i first objects with weight bounded by j . That is,

$$A(i,j) = \max \sum_{k=1}^i x_k v_k \text{ subject to } \sum_{k=1}^i x_k w_k \leq j$$

where each x_k is a non-negative integer.

Assuming $A(0,j) = 0$ for $j = 0 \dots W$, we can fill the table row by row using the following recurrence relation:

$$A(i,j) = \max_{x_i=0 \dots j/w_i} x_i v_i + A(i-1, j - x_i w_i)$$

We only need to consider values of x_i in the range $0 \dots j/w_i$ since any higher value would make $j - x_i w_i$ negative. The final solution will be in $A(n, W)$.

The size of the table is $n \times W$. The time to fill an entry is $O(W)$. Therefore the total time is $O(n \times W^2)$ which is polynomial in $n + W$.

17. [This solution is courtesy of Brian Wongchaowart.]
Let A be an array of size $(n+1) \times (L+1) \times (L+1)$ indexed from 0 and assign $A[i][j][k]$ true if there is a subset S of v_1, \dots, v_i such that

$$\sum_{v \in S} v^3 = j \text{ and } \prod_{v \in S} v = k.$$

```

for (j = 0; j <= L; j++)
    for (k = 0; k <= L; k++)
        A[0][j][k] = false;
A[0][0][1] = true;
for (i = 1; i <= n; i++) {
    for (j = 0; j <= L; j++) {
        for (k = 0; k <= L; k++) {
            A[i][j][k] = A[i - 1][j][k];
            if ((j - v[i]^3 >= 0) && (k % v[i] == 0)
                && A[i - 1][j - v[i]^3][k / v[i]])
                A[i][j][k] = true;
        }
    }
}

```

If $A[n][j][k]$ is true for some $j = k$, then a solution exists and can be reconstructed from the table as follows. (Below, $x[i]$ is assigned 1 if v_i is included in S , and 0 otherwise.)

```

for (j = 0; j <= L; j++)
    if (A[n][j][j])
        break; // solution found
if (j <= L) {
    k = j;
    for (i = n; i >= 1; i--) {
        if ((j - v[i]^3 >= 0) && (k % v[i] == 0)
            && A[i - 1][j - v[i]^3][k / v[i]]) {
            j = j - v[i]^3;
            k = k / v[i];
            x[i] = 1; // include v[i] in the subset
        } else {
            x[i] = 0; // v[i] not included
        }
    }
}

```

Since the size of array A is $O(nL^2)$, the algorithm runs in $O(nL^2)$ time.

18. We present a dynamic programming algorithm for the problem of determining if there are two subsets of n rubies and emeralds that contain the same number of rubies, the same number of emeralds, and the same total value. Let n_e and n_r be the number of emeralds and rubies among the $n = n_e + n_r$ gems. If any of n_e , n_r or L is odd then clearly the problem has no solution. Otherwise we can determine whether there is a solution using dynamic programming.

For $i = 0 \dots n$, $j = 0 \dots n_e/2$, $k = 0 \dots n_r/2$, and $\ell = 0 \dots L$ let $A(i, j, k, \ell)$ be a boolean variable that indicate whether or not there is a subset of the first i gems that contains exactly j emeralds and k rubies and has value ℓ . The table is initialized to all **False**, except for $A(0, 0, 0, 0)$ which is initialized to **True**. We can fill the table row by row using the following code:

```

for  $i = 1 \dots n$  do
  for  $j = 1 \dots n_e/2$  do
    for  $k = 1 \dots n_r/2$  do
      for  $\ell = 0 \dots L/2$  do
        if  $A(i-1, j, k, \ell)$  then
           $A(i, j, k, \ell) = \text{True}$ 
        if  $\text{gem}_i = \text{emerald}$  then
           $A(i, j+1, k, \ell+v_i) \leftarrow \text{True}$ 
        else
           $A(i, j, k+1, \ell+v_i) \leftarrow \text{True}$ 

```

There is a solution to the problem if at the end $A(n, n_e/2, n_r/2, L/2)$ is **True**. The time taken by the algorithm is $O(n^3 \times L)$ which is polynomial in $n + L$.

19. Solution for the word-layout problem:

- (a) Each word can either be placed at the end of the current last line or at the beginning of a new line. Therefore, given a layout of the first i words, there are 2 places to put the $i + 1$ th word. However, note that the first word can only go at the beginning of the first line. Since there are exactly two options for the placement of each word after the first, there are 2^{n-1} total solutions.
- (b) The i th level of the tree contains all possible layouts of the first i words. The left child of a node v corresponds to the layout achieved by appending the $i + 1$ th word to the last line of the layout in v . The right child of v corresponds to the layout achieved by beginning a new line with the $i + 1$ th word.
- (c) Obviously, we can eliminate all layouts that would place more than L characters on a particular line. More helpfully, if there are two solutions S_1 and S_2 of the first i words such that both S_1 and S_2 have k characters on the last line, we only need to keep the one with the smallest total penalty.
- (d) As pruning rule we remark that if we have two layouts of the first i words that have the same number of characters on the last line we only need to keep the layout with the smaller maximum line penalty.

For $i = 0 \dots n$ and $j = 1 \dots L$, let $A(i, j)$ be the smallest maximum line penalty (not counting the last line) of any layout of the first i words that has j characters on the last line. Let w_i be the length of the i th word. Given $A(i-1, j)$ for $j = 1 \dots L$ we can compute $A(i, j)$ by distinguishing three cases:

- i. if $j < w_i$ then $A(i, j) = \infty$ since no layout that ends with the i th word can have less than w_i characters on the last line.
- ii. if $j > w_i$ then $A(i, j) = A(i-1, j-w_i)$ since in this case the i th word is added on the last line, without starting a new one.
- iii. if $j = w_i$ it means that the i th word starts a new line. The penalty depends on the number of characters on the (previous) last line. You want to make the choice that yields the smallest maximum line penalty:

$$A(i, w_i) = \min_{k=1 \dots L} \{ \max(A(i-1, k), L-k) \}$$

Using the above relations it is easy to fill the table. At the end the optimal maximum line penalty is given by the minimum of $A(n, i)$, $i = 1 \dots L$ if you are not counting the last line's penalty. Since here we are counting the last line's penalty, the optimal maximum line penalty is given by the minimum over $i = 1 \dots L$ of $\max(A(n, i), L-i)$. As usual, the actual layout can be reconstructed by tracing the "path" back in the table starting from this final cell.

20. We use dynamic programming to find a subsequence of maximum aggregate cost (note that there is more than one reasonable way to construct a dynamic programming algorithm here). We use the pruning method. Let level l of the tree consist of all subsequences of T that match the first l letters in the pattern P . The pruning rule is that if you have two solutions on the same level are of the same length and end at the same letter in T , then you can prune the one of lesser cost. So let $A[l, s]$ be the maximum cost of a subsequence of T equal to p_1, \dots, p_l , where the last letter of this subsequence is t_s . We then get the following code

```

for  $l = 1 \dots k-1$  do
  for  $s = 1 \dots n$  do
    if  $A[l, s]$  is defined then
      for  $r = s+1 \dots n$  do
        if  $t_r = p_{l+1}$  then
           $A[l+1, r] = \max(A[l+1, r], A[l, s] + c_r)$ 

```

21. We use dynamic programming to compute the optimal solution to the two taxi cab problem. We first note that when p_i is serviced one of the taxi is in p_i while the other taxi is in one of the location $p_0 \dots p_{i-1}$, where p_0

denotes the origin. Therefore at stage i we only need to keep one solution (the best one) for each possible location $p_0 \dots p_{i-1}$ of the second taxi.

For $i = 1 \dots n$ and $j = 0 \dots i - 1$ let $A(i, j)$ be the minimum cost of a routing that serves the points $p_1 \dots p_i$ and leaves one taxi in p_i and the other in p_j . Given the best solutions for stage i we can compute the solutions for stage $i + 1$ by considering the following two possibilities:

- (a) the taxi that was in p_i is used to serves p_{i+1} and the other taxi remains where it was. That is, for $j = 0 \dots i - 1$,

$$A(i + 1, j) = |p_i p_{i+1}| + A(i, j).$$

- (b) the taxi in p_i remains at p_i , and the other taxi serves p_{i+1} . Only the solution with minimum cost is kept.

$$A(i + 1, i) = \min_{k=0 \dots i-1} |p_k p_{i+1}| + A(i, k),$$

Using the above relations it is easy to fill the table. At the end the minimum routing cost is given by the minimum of $A(n, i)$, $i = 0 \dots n - 1$. The actual routing can be reconstructed by tracing the “path” back in the table starting from the final cell.

Since there are n stages and each stage takes $O(n)$ time, the total time is $O(n^2)$ which is polynomial in n .

22. [This solution is adapted from solutions given by Brian Wongchaowart and Daniel Cole.]

Number the points on the line from left to right, so that x_1 is the leftmost point and x_n is the rightmost point. Without loss of generality, assume that the origin is one of these points. Every feasible path begins at the origin and visits the points one by one. The set of all points visited so far must expand by one point to the left or right at a time until every point has been added. Within this set, the point that was added last may be either the leftmost point or the rightmost point. If the last visited point is to the right of the origin, then this point is the rightmost of the visited points. If the last visited point is to the left of the origin, then this point is the leftmost of the visited points. Thus we can think of a tree of feasible solutions where the up to 2^ℓ solutions at level ℓ are all solutions for visiting exactly ℓ points, and each solution has two children depending on whether the next visited point is to the right or left of the currently visited points.

We now give the pruning rule. Consider two solutions at the same level that have the same last visited point. We know from the reasoning above that these two solutions have visited exactly the same points. The two characteristics that we care about these solutions are the total response times of the points that these two solutions have visited, and the response time of the last visited point (which is the length of the path/solution).

We keep the solution that minimizes the total response time to visit the points so far plus $(n - \ell)$ times the length of the path. If you have two solutions at the same level with the same last point, where the difference in their paths lengths is L , then the term $(n - \ell)L$ represents the savings that the solution with the shorter path will incur for the response times of points that will be added in the future.

The outer loop of the code will iterate over levels. Define the array $Cost[\ell, k]$ be the the total response time of the path that visits exactly ℓ points and ends at p_k that minimizes the total response time plus $n - \ell$ times the path length. Define the array $Length[\ell, k]$ be the the length of of the path that visits exactly ℓ points and ends at p_k that minimizes the total response time plus $n - \ell$ times the path length. The next loop loops over the possible values for k . If $Cost[\ell, k]$ is defined and p_k is to the right of the origin, then the next two points that might be visited are p_{k+1} and $p_{k-\ell}$. If

$$Cost[\ell, k] + d(p_k, p_{k+1}) + Length[\ell, k] + (n - \ell - 1)(d(p_k, p_{k+1}) + Length[\ell, k])$$

is less than

$$Cost[\ell + 1, k + 1] + (n - \ell - 1)Length[\ell + 1, k + 1]$$

then $Cost[\ell + 1, k + 1]$ is updated to $Cost[\ell, k] + d(p_k, p_{k+1}) + Length[\ell, k]$ and $Length[\ell + 1, k + 1]$ is updated to $Length[\ell, k] + d(p_k, p_{k+1})$.

If

$$Cost[\ell, k] + d(p_k, p_{k-\ell}) + Length[\ell, k] + (n - \ell - 1)(d(p_k, p_{k-\ell}) + Length[\ell, k])$$

is less than

$$Cost[\ell + 1, k - \ell] + (n - \ell - 1)Length[\ell + 1, k - \ell]$$

then $Cost[\ell + 1, k - \ell]$ is updated to $Cost[\ell, k] + d(p_k, p_{k-\ell}) + Length[\ell, k]$ and $Length[\ell + 1, k - \ell]$ is updated to $Length[\ell, k] + d(p_k, p_{k-\ell})$. If p_k to the left of the origin (and hence the leftmost visited point) then the next two possible visited points are p_{k-1} and $p_{k+\ell}$, and the code is similar to the above. So this gives use time $O(n^2)$.

23. [This solution is courtesy of Brian Wongchaowart.] Number the trips in increasing order by date, so that $d_1 < \dots < d_n$. Let A be an array of size $(n + 1) \times (n + 1)$ indexed from 0. Assign to $A[i][j]$ the minimum possible total cost of taking the first i trips when the last Bahncard was purchased just before trip j (if $j = 0$, no Bahncard is ever purchased). Assume for simplicity that a Bahncard is always purchased on the day of a trip. It is obvious that buying one on a day between two trips cannot result in a lower total cost.

```

for (j = 0; j <= n; j++)
    A[0][j] = 0;

for (i = 1; i <= n; i++) {
    A[i][0] = A[i - 1][0] + f[i];
    for (j = 1; j <= i; j++) {
        if (i == j) {
            // buy Bahncard
            A[i][j] = INFINITY;
            for (k = 0; k < i; k++) {
                if (A[i - 1][k] + f[i] / 2 + BAHNCARD_COST < A[i][j])
                    A[i][j] = A[i - 1][k] + f[i] / 2 + BAHNCARD_COST;
            }
        } else if (i > j && d[i] - d[j] < ONE_YEAR) {
            // use discounted fare
            A[i][j] = A[i - 1][j] + f[i] / 2;
        } else {
            // use full fare
            A[i][j] = A[i - 1][j] + f[i];
        }
    }
}

```

The optimal cost of all n trips is the minimum entry in row n of A . The dates on which a Bahncard was purchased can be determined from A as below, where $x[i]$ is assigned 1 if a Bahncard was purchased just before (on the day of) trip i .

```

i = 0; // last time Bahncard was bought
for (j = 1; j <= n; j++) {
    if (A[n][j] < A[n][i])
        i = j;
}
while (i > 0) {
    x[i] = 1; // buy Bahncard just before trip i
    for (k = 0; k < i; k++) {
        if (A[i][i] == A[i - 1][k] + f[i] / 2 + BAHNCARD_COST)
            i = k; // last time Bahncard was bought
    }
}

```

The algorithm runs in $O(n^2)$ time, since the “for k ” loop runs only once for each i .

24. [This solution is courtesy of Brian Wongchaowart.] Assume for simplicity that $R_n > 0$ (discard any trailing R_i such that $R_i = 0$). For consistency with previous problems let the first integer in sequence R be numbered R_1 , not R_0 as in the problem statement. Let $k = \min(k, n)$, since more server broadcasts than requests are unnecessary. Let A be an array of size $(n + 1) \times (k + 1)$ indexed from 0. Assign to $A[i][j]$ the minimum total waiting time obtainable for the sub-problem consisting of satisfying requests R_1, \dots, R_i using j server broadcasts, so that the optimal solution to the original problem has a total waiting time of $A[n][k]$.

```

for (j = 1; j <= k; j++)
    A[0][j] = 0;

for (i = 1; i <= n; i++) {
    for (j = 1; j <= k && j <= i; j++) {
        if (j == 1) {
            // broadcast at time i + 1
            A[i][j] = 0;
            for (m = 1; m <= i; m++)
                A[i][j] += R[m] * (i - m + 1);
        } else {
            A[i][j] = INFINITY;
            // find best time for broadcast j - 1
            for (b = j - 1; b < i; b++) {
                temp = A[b][j - 1];
                // also broadcast at time i + 1
                for (m = b + 1; m <= i; m++)
                    temp += R[m] * (i - m + 1);
                if (temp < A[i][j])
                    A[i][j] = temp;
            }
        }
    }
}

```

The optimal broadcast times can be determined from A as follows, where $x[i]$ is assigned 1 if there is a broadcast at time i .

```

j = k; // number of broadcasts remaining
i = n;
while (j > 0) {
    x[i + 1] = 1; // broadcast j
    if (j > 1) {
        // find time of broadcast j - 1

```

```

    for (b = j - 1; b < i; b++) {
        temp = A[b][j - 1];
        for (m = b + 1; m <= i; m++)
            temp += R[m] * (i - m + 1);
        if (temp == A[i][j])
            i = b;
    }
}
j--;
}

```

An examination of the loops shows that the algorithm runs in $O(n^2k)$ time.

25. No solution given.
26. [This solution is courtesy of Brian Wongchaowart.] Let A be an array of size $n \times n$ indexed from 1 with every entry initialized to 0. Assign to $A[i][j]$ the number of possible orderings of i objects under the relations $<$ and $=$ where the i objects have been partitioned into j sets of equivalent objects by the $=$ relation, and these equivalent subsets have been ordered by the $<$ relation.

Consider the number of ways in which the addition of one new object to such an ordering of $i - 1$ original objects (making i objects in total) can result in a partition with j equivalent subsets. If $j = 1$, then all of the objects are equivalent and there is never more than one way of doing this. If $j = i$, then all of the objects are different and there are $i!$ ways of ordering them under the $<$ relation.

If $1 < j < i$, then the new object may be equivalent to one of the original objects, so that the number of equivalent subsets stays the same at j , but the new object is added to one of the j existing subsets, requiring a choice between them. It is also possible that the new object is not equivalent to any of the $i - 1$ original objects, in which case there must have been $j - 1$ equivalent subsets of these original objects, and there are now j ways in which the new object can be ordered relative to them under the $<$ relation. Thus array A can be filled in by the following code:

```

A[1][1] = 1;
for (i = 2; i <= n; i++) {
    for (j = 1; j < i; j++)
        A[i][j] = A[i - 1][j] * j;
    for (j = 2; j <= i; j++)
        A[i][j] += A[i - 1][j - 1] * j;
}

```

The sum of the entries in row n of A is the total number of possible orderings of the n objects under the relations $<$ and $=$. Since the size of array A is n^2 , the algorithm runs in $O(n^2)$ time.

27. No solution given.