

AI for Games

Doron Nussbaum

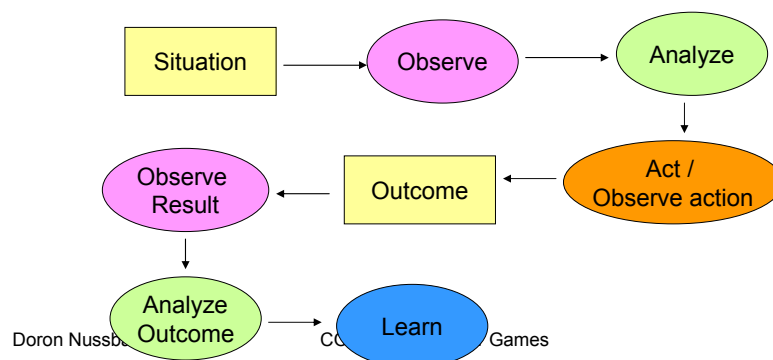
- Introduction
- Movement and path planning
- Games and trees
 - MinMax
 - Decision trees
- Finite State Machines
- Agents
- Other

Intelligence Definition

- “The ability to acquire and apply knowledge and skills” (Oxford dictionary)
- Ability to understand, reason, grasp issues and complex problems/ideas as well as ability to learn from experience of self and others
- **Artificial Intelligence**
 - An attempt by machines (computers) to be intelligent

Humans and Intelligence

- We have the ability to:
 - Learn → look at a situation



AI


- In Academia – understanding intelligence (theory)
 - Thought process
 - Behaviour prediction
 - Comprehension
 - Machine vision
 - Natural language –
 - understanding semantics
 - Carry out a conversation
- Engineering – Use AI to solve real problems
 - Searching (Google)
 - Stock market predictions

Are we there yet?

- Not Yet!
- Large number of applications
 - Machine vision
 - Speech recognition
 - Stock market prediction
 - ...
- Spectrum of domain is hard to achieve
- Some specific areas are promising
 - Chess
 - Carrying out a conversation on specific topics (e.g., whether, daily activities)
- What is in the way
 - Unscripted actions
 - Not always logical
 - Emotions

What about Game AI?

- Game AI is different
 - Provide entertainment
 - Attract the player
 - Should be realistic
- PACMAN
 - Is it AI?
- AI in games provides
 - Interaction with the player
 - Level of unpredictability
 - Somewhat no repetition in the game



http://free-extras.com/images/pacman_game-1973.htm

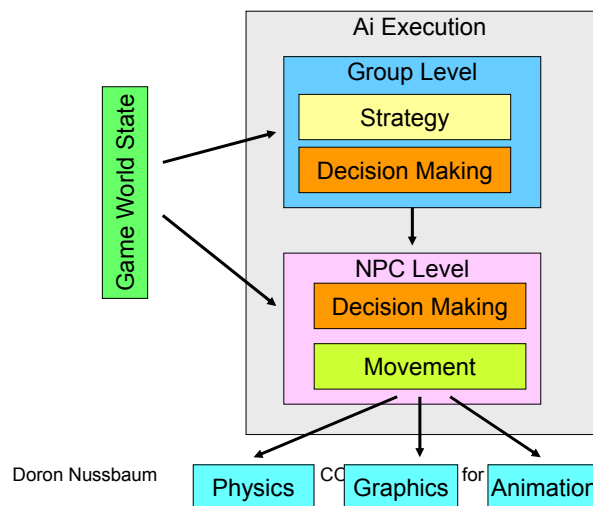
What Should Game AI Be

- Should be good
 - Cannot be too smart – should have some built in flaws
 - Provide fun
- Provide good perception (no obvious flaws)
 - Should not look dumb
 - No unintended flaws – cannot be defeated using a “secret path”
- Must be fast (real time)
- If possible configurable
 - Not hard coded by programmer
- Can adjust to different level of players

What about Game AI?

- Must be smart, but purposely flawed
 - Loose in a fun, challenging way
- No unintended weaknesses
 - No “golden path” to defeat
 - Must not look dumb
- Must perform in real time (CPU)
- Configurable by designers
 - Not hard coded by programmer
- “Amount” and type of AI for game can vary
 - RTS needs global strategy, FPS needs modeling of individual units at “footstep” level
 - RTS most demanding: 3 full-time AI programmers
 - Puzzle, street fighting: 1 part-time AI programmer
 - All of project 2. ☺

AI Model



Movement

- Movement relates to the way an NPC moves in the world
 - Wander
 - Seek
 - Chase / home in / zoom to target
 - Flee

Wander

- NPC moves in the world
 - Without a particular goal
 - Usually aimlessly (no logic in movement)
 - Scouts an area (not exploring)
 - Guarding
 - Cleaning



wander

- Address motion
- Address orientation
- Set up a target
 - Set up a path
 - Path can be a sequence of short line segments.
- Update the motion
 - Speed
 - Velocity
 - Orientation
- Use regular motion equations (distance, speed, velocity)

Seek

- Similar to wander (something in mind)
- Searching for something
 - Treasure
 - Enemy
 - Weapon
- Create constraints – (when is the target visible?)
 - Distance constraints
 - Visibility – colour, size, text
 - Type – searching for a box, target is a sphere

Seek

- Address motion
- Address orientation
- Set up a search pattern
 - Wander
 - Logical scan – in circles, side to side, moving in a maze
- Update the motion
 - Speed
 - Velocity
 - Orientation

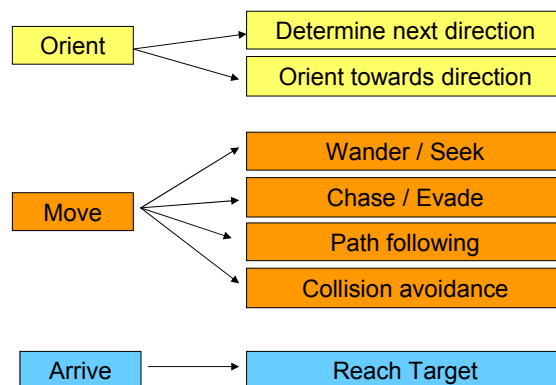
Chase

- Target is visible
 - Known location of target
 - Target may be stationary or in motion
- Realistic actions
 - Change of direction is affected by speed
 - Cannot change direction on the spot (e.g., a car)
- Issues
 - Overshooting the target

Chase

- Address motion
- Address orientation
- Set up a static path
 - Path can be direct
 - Path needs to avoid obstacles
- Set up a chase path (dynamic path)
 - Follow the path
 - Zoom in on current location
 - Prediction of future location
- Update the motion
 - Speed
 - Velocity
 - Orientation

Steering



Path Planning

- Determine how to move in space
 - Wander – move aimlessly
 - Seek – move to a particular location
- What does one have
 - Start point
 - Target point
 - Obstacles
 -

Path Planning

- What is missing?
- Free space!!!
 - Space in which one can move freely
- This may not be trivial
 - Object has area/volume (2D/3D)

Path Planning

- Motion –
 - Assume that object is a point →
 - Move a point in the free space
- Creating free space
 - Convert object to a point
 - Enlarge obstacles accordingly (e.g., Minkowski sums)

Path Planning (last)

- How to move in free space?
 - Hard – not knowing where to go, when to turn
- Visibility
 - Move in a shortest path “notion”
- Attempt to convert the space into a graph

Path Planning

- Types of problems that may be of interest
- Guards placement
 - How many guards are needed to guard area
 - Where to position guards so that the guarded area is covered
- Guarding path
 - Is there a path that a guard can see the guarded area
 - All the time
 - At least once during the motion
 - How many guards are needed?
 - What should the paths be?
- Gaming –
 - Each guard is an autonomous object
 - Place less guards than needed – give the player a chance

Path Planning

- Most of the time path planning is to reach a goal.
 - Shortest path
 -
- What is the meaning of shortest path?

Path Planning

- Input:
 - a graph
 - Vertices
 - Edges
 - Weights
 - Start point
 - End point (in most cases)
- Output: a path (possible \emptyset)
- Algorithm?
 - Path traversing algorithms?

Path traversing Algorithms

- Breadth First Search (BFS)
 - Explore closest neighbourhood first
- Depth First Search (DFS)
 - Explore furthest neighbourhood first

Path Planning

- Algorithms for path planning
 - Best first
 - Dijkstra shortest path
 - A*
 - Hierarchical shortest path

Constructing Graphs

- Depending on the world
 - Grid (cell based)
 - Four connected
 - Eight connected (3x3)
 - 16 connected (5x5)
 - Vector based
 - TIN
 - Obstacles
 - Free space

Best First

- Usually used on a grid based graphs
- Idea
 - Attempt to move as “fast” as possible to the target (Greedy algorithm)
 - Attraction to the target/goal position

Dijkstra Shortest Path

- Search for the target around the start point until target is found.
 - No relation to the target point
- Algorithm properties
 - Triangle inequality
$$\text{cost}(\pi(u,w)) \leq \text{cost}(\pi(u,v)) + \text{cost}(\pi(v,w))$$
 - $\delta(u)$ is the minimum $\text{cost}(\pi(u,v))$

Dijkstra Shortest Path

- Let (v,u) be an edge in G
 - $\delta(u) \leq \delta(v) + \text{weight}(u,v)$
- If v_0, v_1, \dots, v_k be a shortest path from v_0 to v_k then v_i, \dots, v_j is a shortest path from v_i to v_j where $0 \leq i, j \leq k$ and $i < j$

- $v_i \leftarrow \infty$
- $s \leftarrow 0$
- Insert all vertices to a priority queue Q
- While ($Q \neq \emptyset$)
 - $u \leftarrow \text{top}(Q)$
 - for all v which are neighbours of u
 - if $\text{cost}(\pi(s,u)) + \text{weight}(u,v) < \text{cost}(\pi(s,v))$ then
 - $\text{cost}(\pi(s,v)) \leftarrow \text{cost}(\pi(s,u)) + \text{weight}(u,v)$
 - Update parent of v
 - Update Q with new cost of v

What is the “problem” with Dijkstra?

What is the “problem” with Dijkstra?

- A BFS algorithm
- Search everywhere without any relationship to the target
- Solution ?

What is the “problem” with Dijkstra?

- A BFS algorithm
- Search everywhere without any relationship to the target
- Solution
 - Combine Best First and Dijkstra

A* Path Algorithm

- A heuristic algorithm
- Attempts to take the best of both worlds
 - The BFS behaviour of Dijkstra
 - The DFS behaviour of Best First

A* Path Algorithm

- Modify the comparison statement of Dijkstra priority queue algorithm
- Instead of extracting u such that $\text{cost}(\pi(s,u))$ is the minimum in Q
- Use $\text{cost}(\pi(s,u)) + \text{Estimate}(u,t)$

A* Path Algorithm

- What kind of estimate one can use?

Large Graphs


- What can be done with large graphs?
- How can many queries be handled?
- Scalability
 - Domain size
 - Number of players
- Solution
 - Speed up queries
 - Merge queries




Hierarchical Graphs



- Create a hierarchy
- Similar to a highway system
 - Neighbourhood roads
 - City roads
 - Regional roads


Graph Hierarchy

- Designate some paths as “high” level paths (highways)
 - Few connections
 - Few edges
- Create a representation of the space
 - Possibly a sequence of representations $G, G^1, G^2 \dots$
- Operation
 - Gravitate to a high level representation
 - Find path on high level
 - Convert path back to a low level “real path”









Game Playing Programs

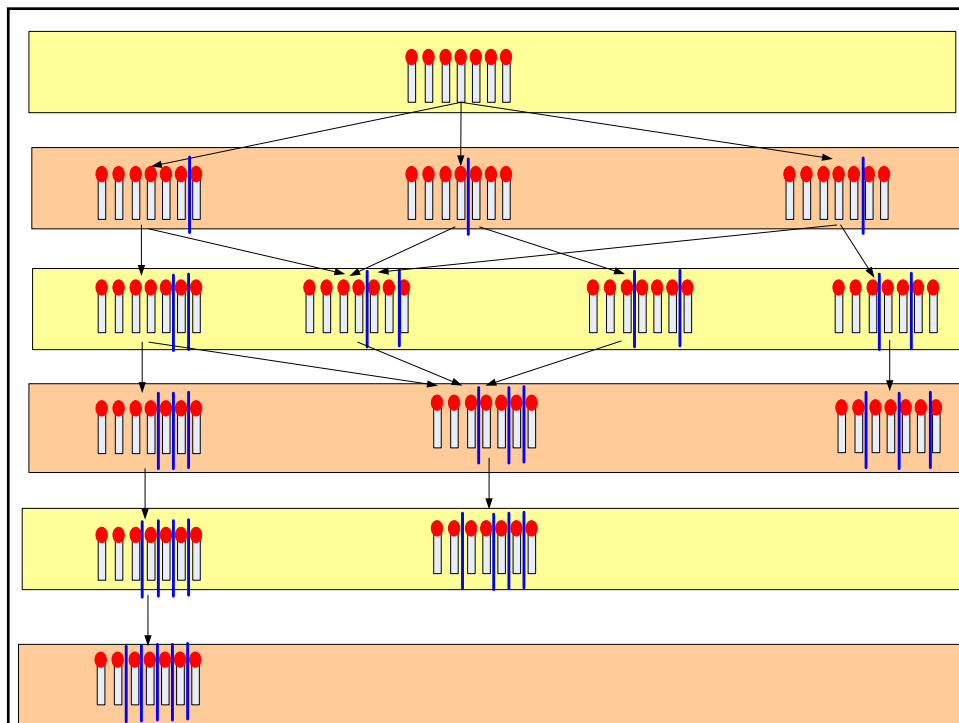

- Game playing programs obey some rules:
 - Visibility – is the board visible/partial visible
 - Turns – is the game played in order
 - Chance/Luck – what is the nature of a “move”
 - Governed by probability (card game, dice)

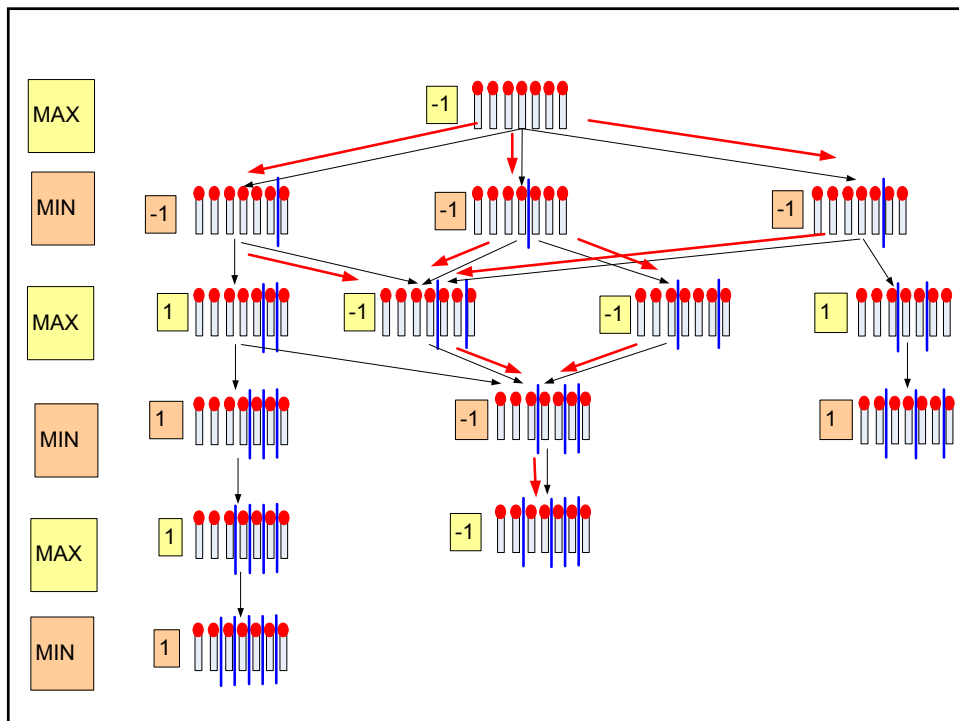
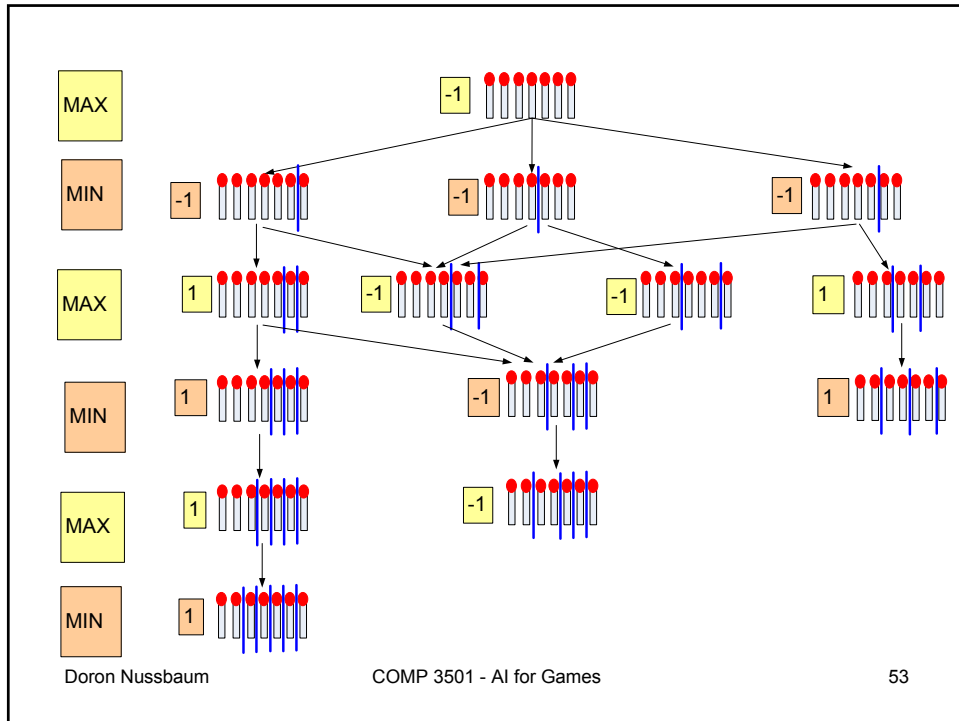
How to play?

- Deterministic game
- Create a tree with all possible moves
- Search the tree for the best move
 - DFS
- Assuming that each player plays for the best move then each try to follow a path that guarantees a “victory”

Division Nim Game

- Given a pile of matches
- Each player in his/her turn takes one pile and divides it into two piles of different size
- Game ends when there is non other option to play
- For example given 6 matches one can divide into
 - 1-5
 - 2-4





$$\text{minimax}(v) = \begin{cases} \text{value}(v) & \text{if } v \text{ is a leaf} \\ \min_{u \in \text{children}(v)} \{\text{minimax}(u)\} & \text{if } v \text{ is a MIN node} \\ \max_{u \in \text{children}(v)} \{\text{minimax}(u)\} & \text{if } v \text{ is a MAX node} \end{cases}$$

In: node v in the tree
 Out: value of v

```

MiniMax(v)
if (children(v) = ∅) then
  return value(v)
else if label(v) = MIN then
  d ← +∞
  for all u ∈ children(v) do
    d ← min{d, MiniMax(u)}
  endfor
  return d
else
  d ← -∞
  for all u ∈ children(v) do
    d ← max{d, MiniMax(u)}
  endfor
  return d
endif
  
```

Complexity

- Assuming a branching factor is b and a depth of the search tree is k then the tree size is

$$1 + b + b^2 + \dots + b^k = \frac{1 - b^{k+1}}{1 - b} = \frac{b^{k+1} - 1}{b - 1}$$

- Need to traverse the tree (DFS/post order)
- $O(b^k)$

MiniMax Trees

- Complete? Yes (if tree is finite)
- Optimal? Yes (against an optimal opponent)
- Time complexity $O(b^k)$
- Space complexity $O(b^k)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible

Partial Trees

- At times it is not possible to build a full tree
 - Time complexity $O(b^k)$
 - Space complexity $O(b^k)$
- For example - Chess game
 - Branching factor - $b \approx 35$,
 - A reasonable game is $k \approx 100$
 - Exact solution completely infeasible
 - $35^{100} \approx 2^{500}$

Evaluation Function





- Solution
 - Build a partial tree
- Question –
 - What value to give a node in the “middle” of the tree?
- Develop an evaluation function to assess the state of the game

Searching

- if search **depth limit** was reached
 - **Evaluate/Compute** state of current position
 - Return value
- else
 - if MAX level then
 - apply MiniMax to each child
 - return maximum of results
 - else // MIN level
 - apply MiniMax to each child
 - return minimum of results

Evaluate Tic-Tac-Toe

- Evaluate the number of free moves that are available to win




- $\text{Eval}(p)$ –
number of directions for Max - number of
directions for Min
- $\text{Eval}(p) = \infty$ if Max wins
- $\text{Eval}(p) = -\infty$ if Min wins

- $\text{Eval}(p) = 6 - 5 = 1$

X	O	


Example – Assume 2 steps ahead

- Max Move?




Example – Assume 2 steps ahead

- Min Move




Example – Assume 2 steps ahead

- Min Move



Example – Assume 2 steps ahead

- Min Move




Doron Nussbaum

COMP 3501 - AI for Games

67

Example – Assume 2 steps ahead

- Min Move




Doron Nussbaum

COMP 3501 - AI for Games


68

Example – Assume 2 steps ahead

- Min Move




- Max move




Pruning the tree

- Tree can be quite large
 - Evaluation consumes computing cycles
- Trimming the tree → saves work
- Alpha-beta pruning (α - β pruning)


Example




Example




Example



Example




Example




Game Agents

- There are several ways of representing non-player-characters (NPC)
 - Game agents is one option which is easy to relate to
 - In most cases it is an enemy, or an ally
 - Sometimes it is neutral (background NPC)
- Agents operate through a sense-think-act cycle



Game Agents

- There are several ways of representing non-player-characters (NPC)
 - Game agents is one option which is easy to relate to
 - In most cases it is an enemy, or an ally
 - Sometimes it is neutral (background NPC)
- Agents operate through a sense-think-act-learn cycle



Sensing

- Agents use sensing to gather world information:
 - Barriers (obstacles)
 - Environment (forest, lake, urban setting)
 - Opponents (position and state)
 - Other objects/agents
- Sensing acts as a stimulator to the agent which may trigger an action
- Agents can query the game for information (cheating)
 - Impose similar constraints on data as player
 - Cannot see through walls
 - Sensing has limitation (hearing, vision)
 - Can only use what they have experienced

Sensing - Vision

- A fundamental sense in gaming
 - An agent may ask to get a list of all enemies
 - Not all enemies are known to agent at that point
 - Time consuming to provide a list of all
 - Not all enemies are relevant at this point
 - Limit knowledge to what is visible
- Visibility is compute intensive and can be complicated
 - Partial visibility of objects
 - Complex terrain

Sensing - Vision

- Game must be able to answer queries fast
 - Is the object within viewing range of agent?
 - Is the object within viewing angle of agent?
 - Is the object unobstructed?
- Compute a vector to each object
 - Check magnitude
 - Check angle (dot-product) to determine field of view
 - Check if obscured

Sensing - Hearing

- Player crawls (tip-toe) past an enemy → enemy does not hear
 - Run by enemy → enemy hears
 - Player suddenly shoots a gun at night → agents may hear it and start to move towards sound (e.g., urban war)
- Implement as event driven
 - When player performs action then notify agents within range
 - Obstacles (compute intensive)
 - Variations –
 - Other sounds muffle player
 - Agents are listening → increase range of sound

Sensing - Communication

- Agents can communicate with each other
 - Transferring sensed knowledge to other agents
 - Instant – using a two-way radio
 - Slow – agents runs to other agents and informs
 - Via sound – shouting
- Agents senses data from other agents
 - Vision – signs, colours, actions, not seing
 - Sound
- Reaction time
 - Sensing may take time (e.g., instant reaction to alarm)
 - Build in an artificial delay (e.g., using timer)

Thinking

- Evaluate the gathered information
 - The crux of the AI system
- Can be complex or simple (as required)
- Generally two ways
 1. Pre coded expert knowledge
 - Typically handcrafted - a sequence of “if-then” rules with added randomness
 2. Search algorithm for optimal/best solution
 1. MinMax trees

Thinking

- Expert Knowledge (Finite State Machine, Decision Trees)
- Appealing
 - Simple to create
 - Natural (can think about own behaviour)
 - If “I am tired” then “find a place to sleep”
 - Embodies common sense and knowledge of domain
 - Is enemy weaker than me? → attack
 - Is enemy stronger than me? → run away and/or seek help
- Often not scalable
 - Can be complex with many rules
 - Changing one rule may affect others
- Still sufficient as most agents address a limited domain
 - carry out very limited functionality

Thinking

- Search
 - Look ahead and see what to do next
 - Path finding (A*)
 - Next move (MinMax)
 - Works well with existing knowledge
 - Visible obstacles, current location
- Machine learning
 - Evaluate past actions → use for future actions
 - Techniques are improving but
 - Compute intensive
 - Domain specific

Acting

- Sensing and thinking are invisible to player
 - No matter how sophisticated the these two steps the player is oblivious to them
- Acting is visible to user
 - Game must demonstrate its sophistication via realizable actions
- Common actions
 - Pickup weapon
 - Fire
 - Communicate with player
 - animation

Acting

- Learning and remembering
 - May not be important when agent is short lived
 - If agent is alive for more than 30s can be helpful
 - If player always attacks from left then shield to the left
- Implementation does not have to be complex
 - Gather stats on player behaviour (direction of attack)
 - Information can fade (limited capacity for past)
 - Simple things such as using last player's location to start a search
 - Can be done at the game level –
 - Too many agents are killed in an area → create a smell, colour the area in blood, leave part of agents on the ground (armoury, part of weapons)

Acting

- Making agents stupid
 - Many cases agents can be superior and dominate the player
 - Agent always makes headshot or takes cover
 - Agents runs faster
 - Can be considered as cheating (inside trader)
 - Dumb the agent by giving “human” traits
 - Less accuracy
 - Longer reaction time
 - Engaging the player one at a time
 - Make mistakes
 - Make unnecessary moves
 - Run out of ammunition at critical time → force to retreat

Acting

- Agent Cheating
 - Ideally agents should not be superior by having an unfair advantage
 - Better attributes
 - More attributes
 - Inside knowledge
 - Agents may resort to “cheating” at a high level in order to challenge the player.
 - Player should be aware of it → player will be ready
 - May challenge the player
 - Compute intensive reasons
 - Development time
 - Let the player know upfront